

CSC263 Week 2

If you feel rusty with probabilities, please read the Appendix C of the textbook. It is only about 20 pages, and is highly relevant to what we need for CSC263.

Appendix A and B are also worth reading.

This week topic

- ADT: **Priority Queue**
- Data structure: **Heap**

An ADT we already know

First in first serve

Queue:

→ a collection of elements

→ supported operations

◆ Enqueue(Q, x)

◆ Dequeue(Q)

◆ PeekFront(Q)



The new ADT

Max-Priority Queue:

→ a collection of elements **with priorities**, i.e., each element x has $x.\text{priority}$

→ supported operations

◆ **Insert**(Q, x)

- like `enqueue(Q, x)`

◆ **ExtractMax**(Q)

- like `dequeue(Q)`


◆ **Max**(Q)

- like `PeekFront(Q)`

◆ **IncreasePriority**(Q, x, k)

- increase $x.\text{priority}$ to k

Oldest person first



Applications of Priority Queues

- Job scheduling in an operating system
 - ◆ Processes have different priorities (Normal, high...)
- Bandwidth management in a router
 - ◆ Delay sensitive traffic has higher priority
- Find minimum spanning tree of a graph
- etc.

**Now, let's implement
a (Max)-Priority Queue**

40 -> 33 -> 18 -> 65 -> 24 -> 25

Use an **unsorted linked list**

→ **INSERT(Q, x)** # x is a node

◆ Just insert x at the head, which takes $\Theta(1)$

→ **IncreasePriority(Q, x, k)**

◆ Just change x.priority to k, which takes $\Theta(1)$

→ **Max(Q)**

◆ Have to go through the whole list, takes $\Theta(n)$

→ **ExtractMax(Q)**

◆ Go through the whole list to find x with max priority ($O(n)$), then delete it ($O(1)$ if doubly linked) and return it, so overall $\Theta(n)$.

65 -> 40 -> 33 -> 25 -> 24 -> 18

Use a **reversely sorted linked list**

→ **Max(Q)**

- ◆ Just return the head of the list, $\Theta(1)$

→ **ExtractMax(Q)**

- ◆ Just delete and return the head, $\Theta(1)$

→ **INSERT(Q, x)**

- ◆ Have to linearly search the correct location of insertion which takes $\Theta(n)$ in worst case.

→ **IncreasePriority(Q, x, k)**

- ◆ After increase, need to move element to a new location in the list, takes $\Theta(n)$ in worst case.

$\Theta(1)$ is fine, but $\Theta(n)$ is kind-of bad...

unsorted linked list

sorted linked list

...

Can we link these elements in a **smarter way, so that we never need to do $\Theta(n)$?**




Yes, we can!

Worst case running times

	unsorted list	sorted list	Heap
Insert(Q, x)	$\Theta(1)$	$\Theta(n)$	$\Theta(\log n)$
Max(Q)	$\Theta(n)$	$\Theta(1)$	$\Theta(1)$
ExtractMax(Q)	$\Theta(n)$	$\Theta(1)$	$\Theta(\log n)$
IncreasePriority (Q, x, k)	$\Theta(1)$	$\Theta(n)$	$\Theta(\log n)$





Binary Max-Heap



A **binary max-heap** is a **nearly-complete binary** tree that has the **max-heap property**.




It's a **binary** tree

Each node has **at most** 2 children




It's a **nearly-complete** **binary** tree

Each level is **completely filled**, except the bottom level where nodes are filled to as **far left** as possible



Why is it important to be a **nearly-complete binary** tree?

Because then we can **store** the tree in an **array**, and each node knows which **index** has the its parent or left/right child.



A	B	C	D	E	F
index: 1	2	3	4	5	6

$$\text{Left}(i) = 2i$$

$$\text{Right}(i) = 2i + 1$$

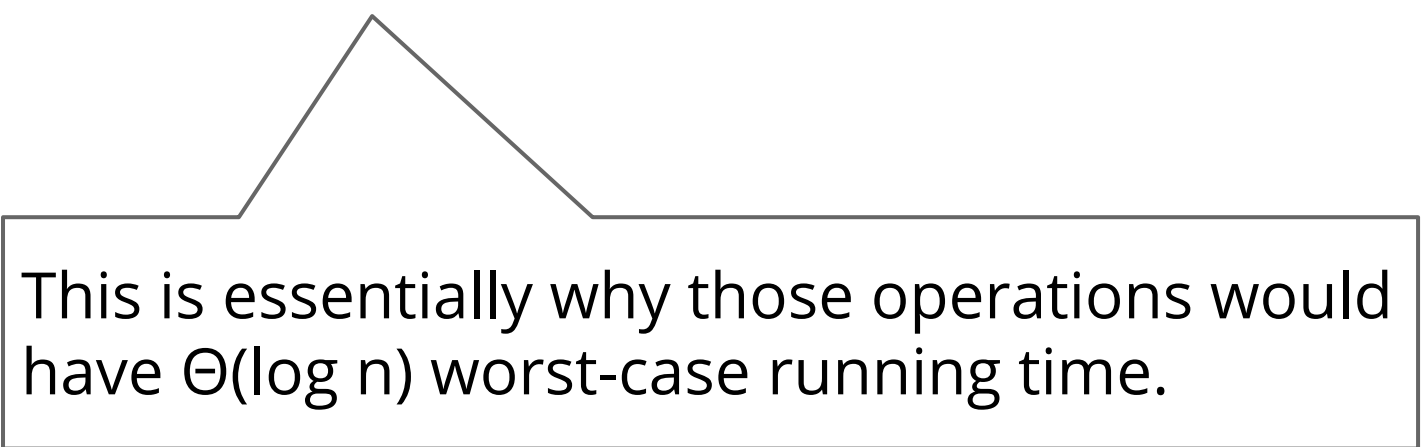
$$\text{Parent}(i) = \text{floor}(i/2)$$

Assume index starts from 1

Why is it important to be a **nearly-complete binary** tree?

Another reason:

The **height** of a complete binary tree with **n** nodes is **$\Theta(\log n)$** .




This is essentially why those operations would have $\Theta(\log n)$ worst-case running time.

A thing to remember...

A heap is stored in an array.




Binary Max-Heap




A **binary max-heap** is a **nearly-complete binary** tree that has the **max-heap property**.

The **max-heap** property

Every node has key (priority) greater than or equal to keys of its **immediate** children.




The **max-heap property**



Every node has key (priority) greater than or equal to keys of its **immediate** children.

Implication: every node is larger than or equal to **all its descendants**, i.e., every **subtree** of a heap is also a heap.



We have a binary max-heap defined,
now let's do operations on it.

- $\text{Max}(Q)$
- $\text{Insert}(Q, x)$
- $\text{ExtractMax}(Q)$
- $\text{IncreasePriority}(Q, x, k)$

Max(Q)

Return the largest key in Q,
in $O(1)$ time


Max(Q): return the maximum element

Return the **root** of the heap, i.e.,

just **return Q[1]**


(index starts from 1)

worst case **$\Theta(1)$**



Q

65	40	25	33	24	18
----	----	----	----	----	----



Insert(Q, x)


Insert node x into heap Q ,
in $O(\log n)$ time

Insert(Q, x): insert a node to a heap

First thing to note:

Which spot to add the new node?

The only spot that keeps it a **complete** binary tree.




Increment heap size

Insert(Q, x): insert a node to a heap

Second thing to note:

Heap property might be broken, how to fix it and **maintain** the heap property?

“**Bubble-up**” the new node to a proper position, by **swapping with parent**.




Insert(Q, x): insert a node to a heap

Second thing to note:

Heap property might be broken, how to fix it and **maintain** the heap property.

“**Bubble-up**” the new node to a proper position, by **swapping with parent**.




Insert(Q, x): insert a node to a heap

Second thing to note:
Heap property might be broken, how to fix it and **maintain** the heap property.

“**Bubble-up**” the new node to a proper position, by **swapping with parent**.

Worst-case:
 $\Theta(\text{height}) = \Theta(\log n)$



ExtractMax(Q)


Delete and return the largest key in Q,
in $O(\log n)$ time

ExtractMax(Q): delete and return the maximum element

First thing to note:

Which **spot** to remove?

The only **spot** that keeps it a **complete** binary tree.



Decrement heap size

ExtractMax(Q): delete and return the maximum element


First thing to note:

THIS guy's key (root) should be deleted.


Which **spot** to remove?

The only **spot** that keeps it a **complete** binary tree.

But the last guy's key should NOT be deleted.



Overwrite root with the **last guy's** key, then **delete** the last guy (decrement heap size).




Decrement heap size


ExtractMax(Q): delete and return the maximum element

Now the **heap property** is broken again..., need to fix it.



“**Bubble-down**” by swapping with...
a child...



Which child to swap with?



so that, after the swap, **max-heap property** is satisfied




The “elder” child!

because it is the **largest** among the three

ExtractMax(Q): delete and return the maximum element

Now the **heap property** is broken again..., need to fix it.


“**Bubble-down**” by swapping with **the elder child**



ExtractMax(Q): delete and return the maximum element

Now the **heap property** is broken again..., need to fix it.


“**Bubble-down**” by swapping with...
the elder child



ExtractMax(Q): delete and return the maximum element

Now the **heap property** is broken again..., need to fix it.

“**Bubble-down**” by swapping with **the elder child**



Worst case running time: $\Theta(\text{height})$ + some constant work
 $\Theta(\log n)$

Quick summary

Insert(Q, x):

→ Bubble-up, swapping with parent

ExtractMax(Q)

→ Bubble-down, swapping elder child

Bubble up/down is also called **percolate** up/down, or **sift** up down, or **tickle** up/down, or **heapify** up/down, or **cascade** up/down.

CSC263 Week 2

Thursday

Announcements

Problem Set 2 is out


→ due next Tuesday 5:59pm

Additional office hours on Mondays

→ 4 - 5:30pm (or by appointment)

A quick review of Monday

→ Max-Priority Queue implementations

- ◆ unsorted and sorted linked list -- $O(1)$, $O(n)$
- ◆ binary max-heap -- $O(1)$, $O(\log n)$
 - $\text{Max}(Q)$
 - $\text{Insert}(Q, x)$
 - bubble up - swapping with parent
 - $\text{ExtractMax}(Q)$
 - bubble down - swapping with elder child
 - $\text{IncreasePriority}(Q, x, k)$ 


IncreasePriority(Q, x, k)

Increases the key of node x to k ,
in $O(\log n)$ time

IncreasePriority(Q, x, k):
increase the key of node x to k

Just increase the key,
then...


Bubble-up by
swapping with parents,
to proper location.



IncreasePriority(Q, x, k):
increase the key of node x to k

Just increase the key,
then...

Bubble-up by
swapping with parents,
to proper location.



Worst case running time: $\Theta(\text{height})$ + some constant work
 $\Theta(\log n)$

Now we have learned how implement a priority queue using a heap

- $\text{Max}(Q)$
- $\text{Insert}(Q, x)$
- $\text{ExtractMax}(Q)$
- $\text{IncreasePriority}(Q, x, k)$


Next:

- How to use heap for **sorting**
- How to **build a heap** from an unsorted array

HeapSort

Sorts an array, in $O(n \log n)$ time

The idea




Worst-case running time: each ExtractMax is **$O(\log n)$** , we do it **n** times, so overall it's...
 $O(n \log n)$

How to get a sorted list out of a heap with n nodes?

Keep extracting max for n times, the keys extracted will be sorted in non-ascending order.

Now let's be more precise

What's needed: modify a max-heap-ordered **array** into a **non-descendingly** sorted **array**



We want to do this “**in-place**” without using any extra array space, i.e., just by **swapping** things around.

Valid heaps are green rectangled

65	40	25	33	18	24
----	----	----	----	----	----

24	40	25	33	18	65
----	----	----	----	----	----

24	40	25	33	18	65
----	----	----	----	----	----

40	33	25	24	18	65
----	----	----	----	----	----

18	33	25	24	40	65
----	----	----	----	----	----

33	25	18	24	40	65
----	----	----	----	----	----

25	24	18	33	40	65
----	----	----	----	----	----

24	18	25	33	40	65
----	----	----	----	----	----

18	24	25	33	40	65
----	----	----	----	----	----

18	24	25	33	40	65
----	----	----	----	----	----

Step 1: swap first (65) and last (24), since the tail is where 65 (max) belongs to.

Step 2: decrement heap size

This node is like deleted from the tree, not touched any more.

Step 3: fix the heap by **bubbling down** 24

Repeat Step 1-3 until the array is fully sorted (at most n iterations).

HeapSort, the pseudo-code

HeapSort(A)

“sort any array A into non-descending order”

Missing!


BuildMaxHeap(A) *# convert any array A into a heap-ordered one*

for $i \leftarrow A.size$ downto 2:

 swap $A[1]$ and $A[i]$ *# Step 1: swap the first and the last*

$A.size \leftarrow A.size - 1$ *# Step 2: decrement size of heap*

 BubbleDown(A, 1) *# Step 3: bubble down the 1st element in A*



Does it work?

It works for an array A that is initially **heap-ordered**, it does work NOT for **any array**!


BuildMaxHeap(A)

Converts an array into a max-heap
ordered array, in $O(n)$ time

Convert any array into a heap ordered one

any array



18	33	25	65	24	40
----	----	----	----	----	----



heap ordered array

65	40	25	33	18	24
----	----	----	----	----	----

In other words...



Idea #1

```
BuildMaxHeap(A):
```

```
    B ← empty array # empty heap
```

```
    for x in A:
```

```
        Insert(B, x) # heap insert
```

```
    A ← B # overwrite A with B
```

Running time:

Each Insert takes **$O(\log n)$** , there are **n** inserts...
so it's **$O(n \log n)$** , not very exciting.

Not **in-place**, needs a second array.


WHAT IF I TOLD YOU

YOU CAN DO BETTER THAN THIS

imgflip.com

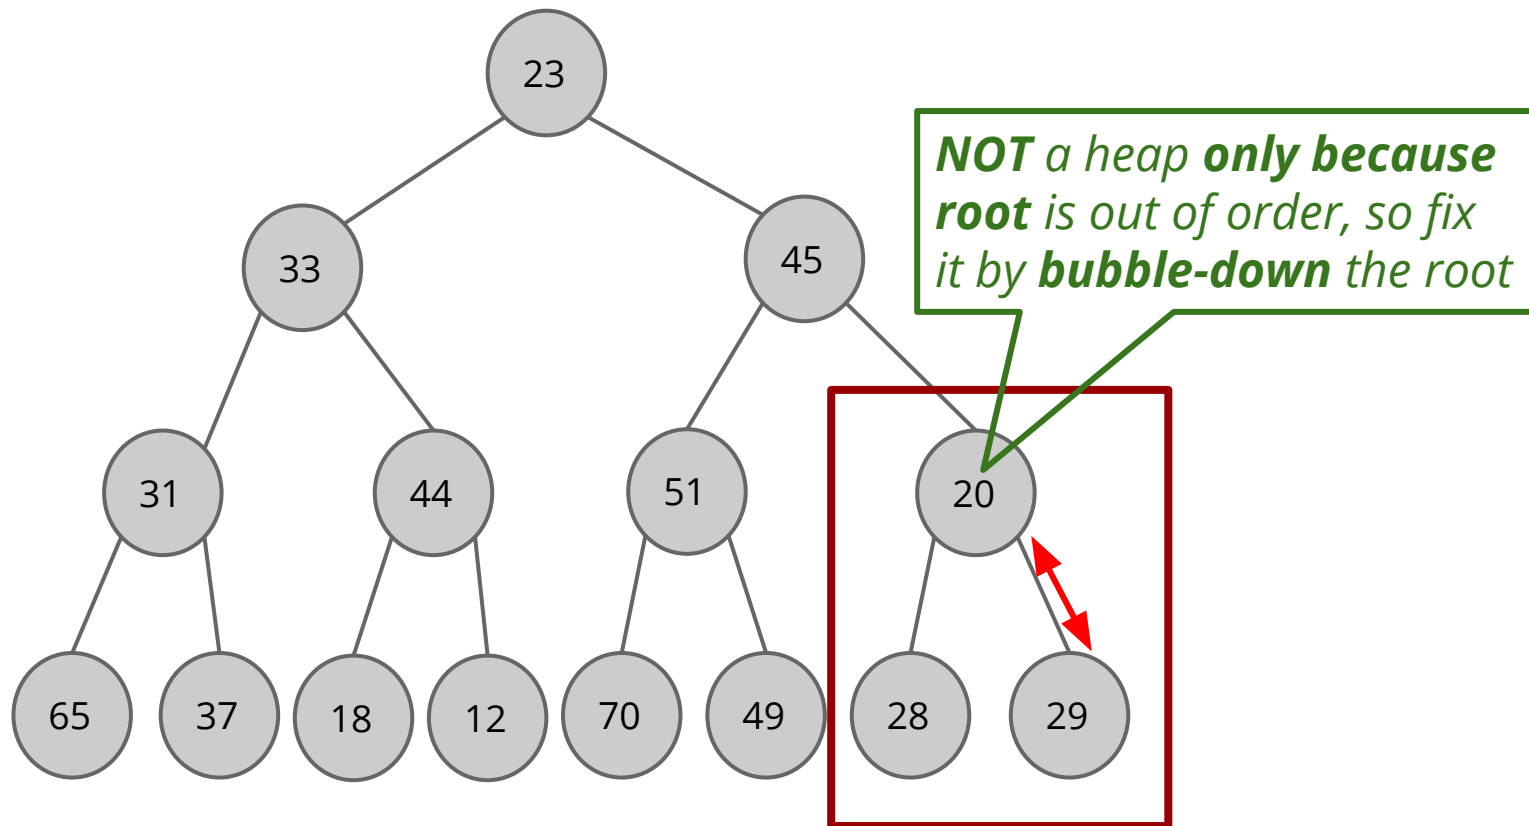
Idea #2

Fix heap order, from bottom up.



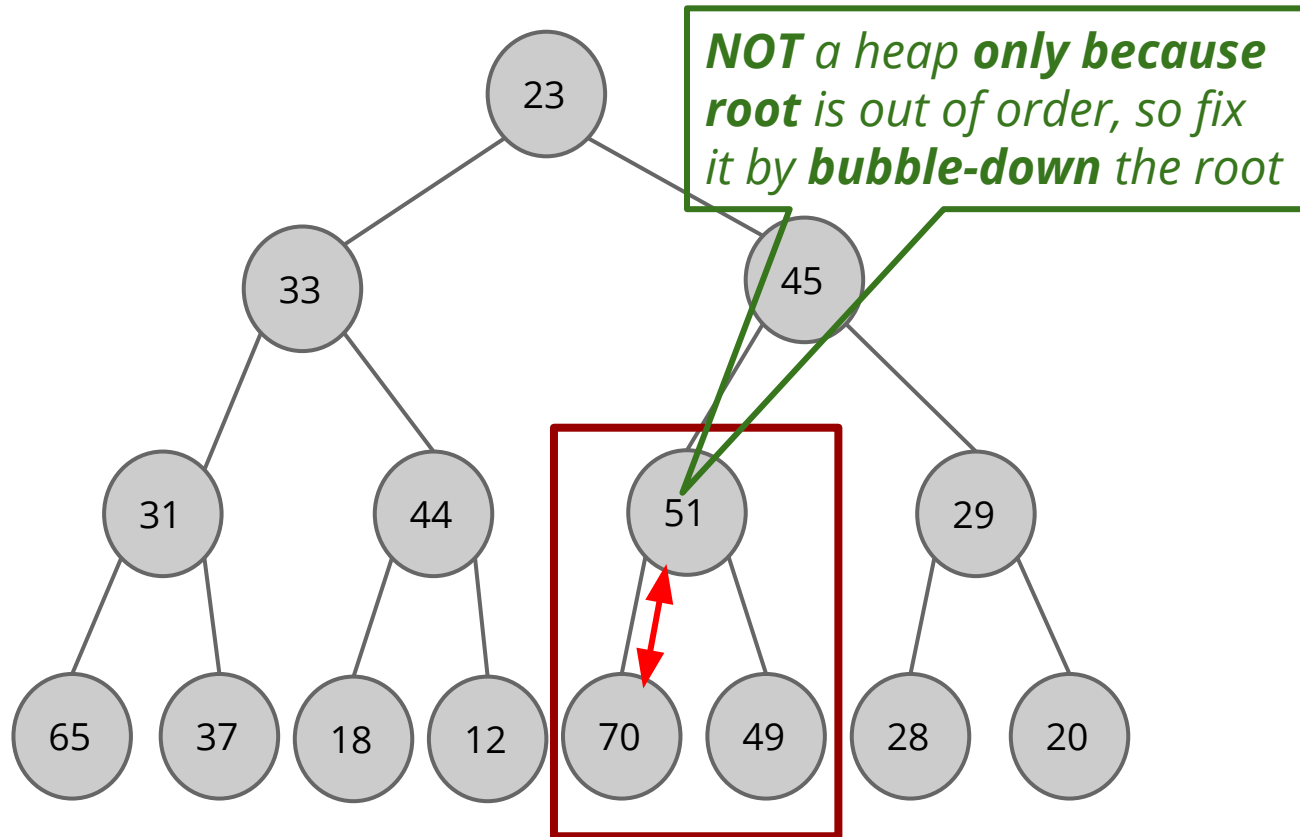
Idea #2

Adjust heap order, from bottom up.




Idea #2

Adjust heap order, from bottom up.



Idea #2

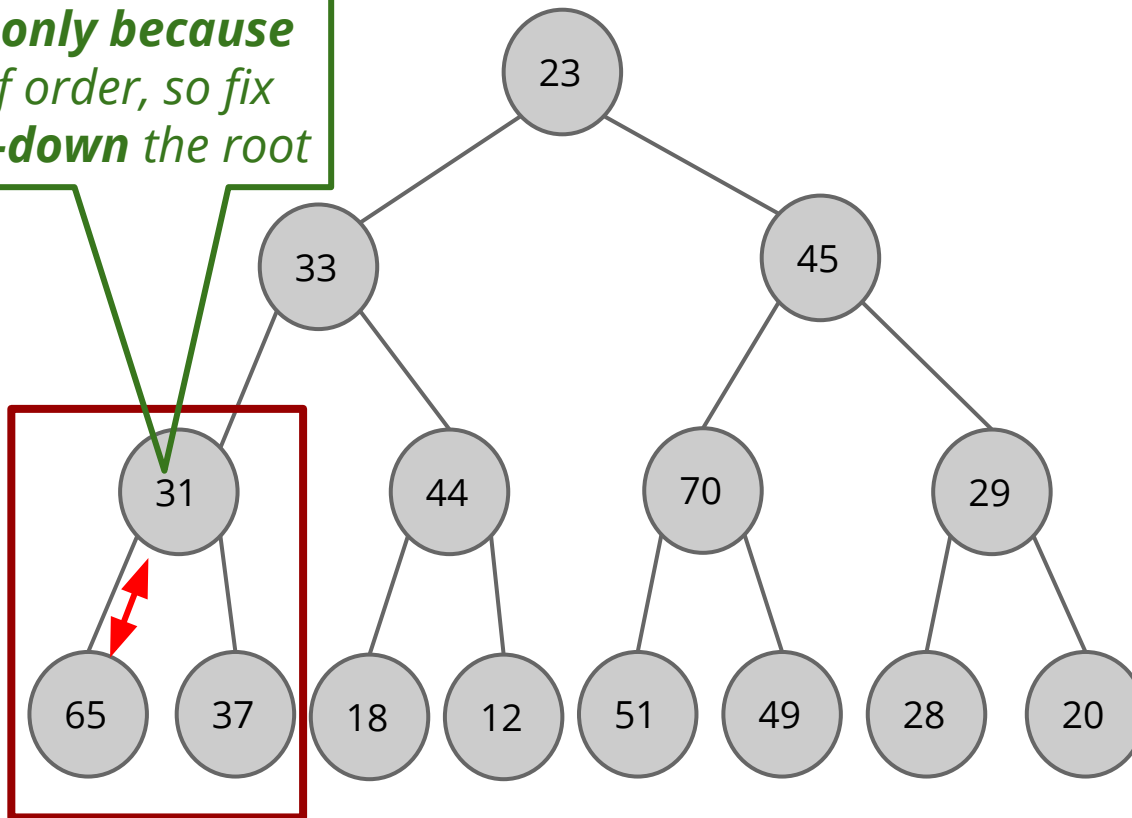
Adjust heap order, from bottom up.



Idea #2

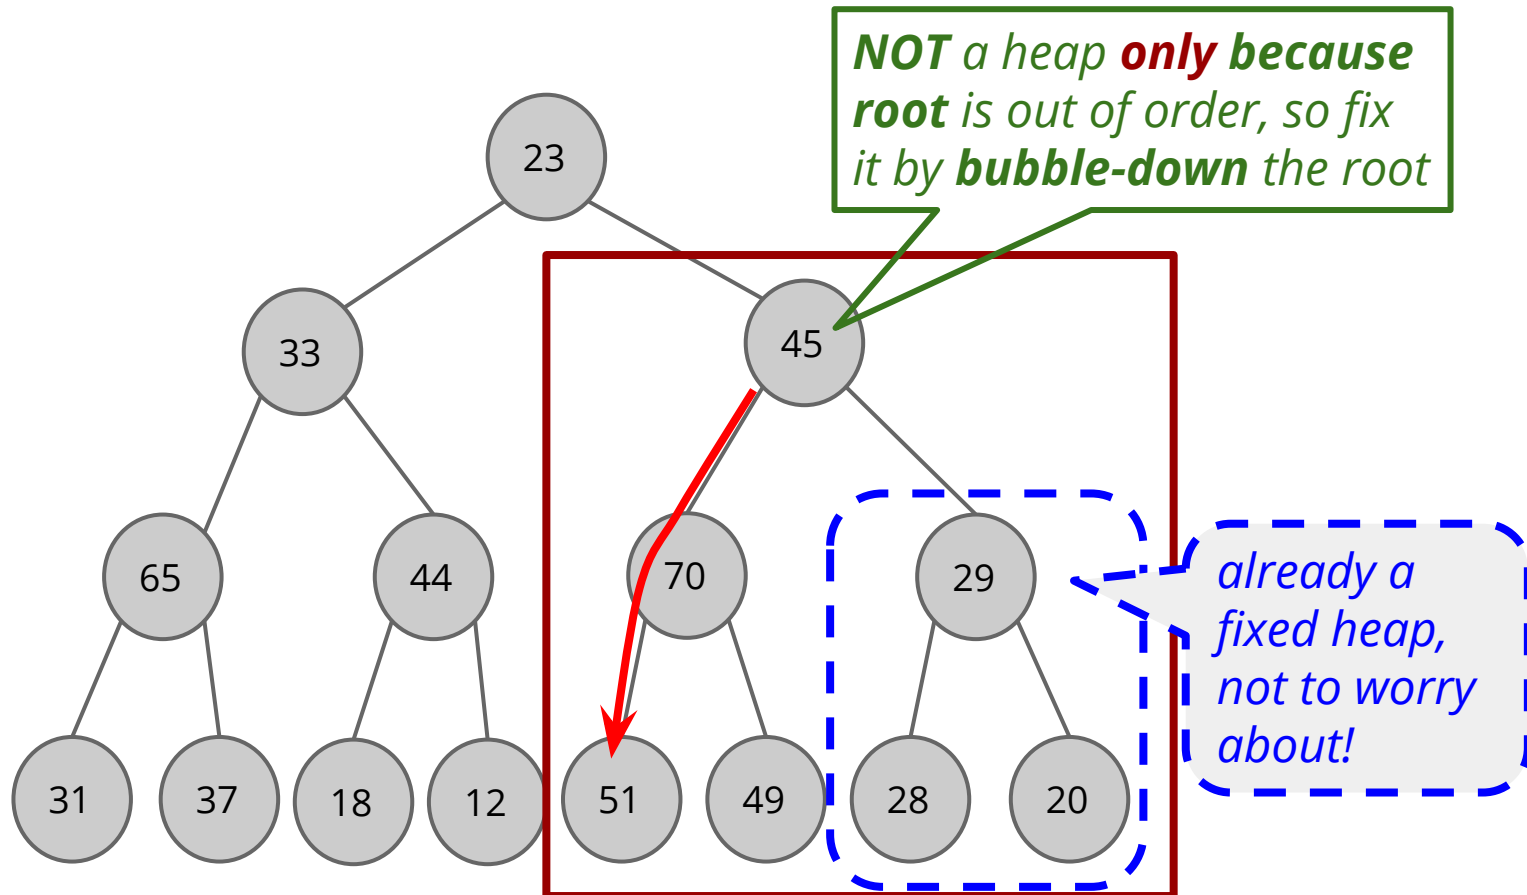
Adjust heap order, from bottom up.

NOT a heap **only because**
root is out of order, so fix
it by **bubble-down** the root



Idea #2


Adjust heap order, from bottom up.



Idea #2

Adjust heap order, from bottom up.

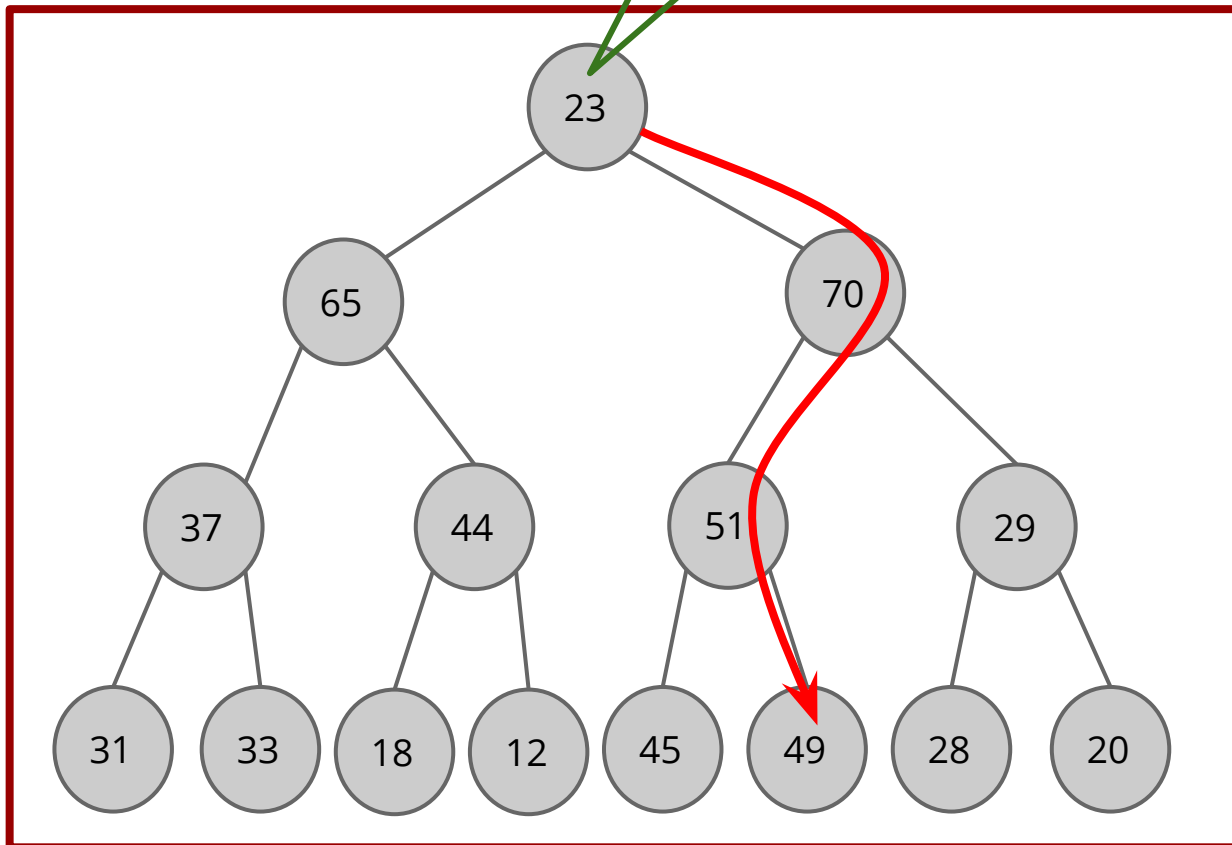
*NOT a heap **only** because root is out of order, so fix it by **bubble-down** the root*



Idea #2

***NOT** a heap **only** because **root** is out of order, so fix it by **bubble-down** the root*


Adjust heap order, from bottom up.




Idea #2

Adjust heap order, from bottom up.

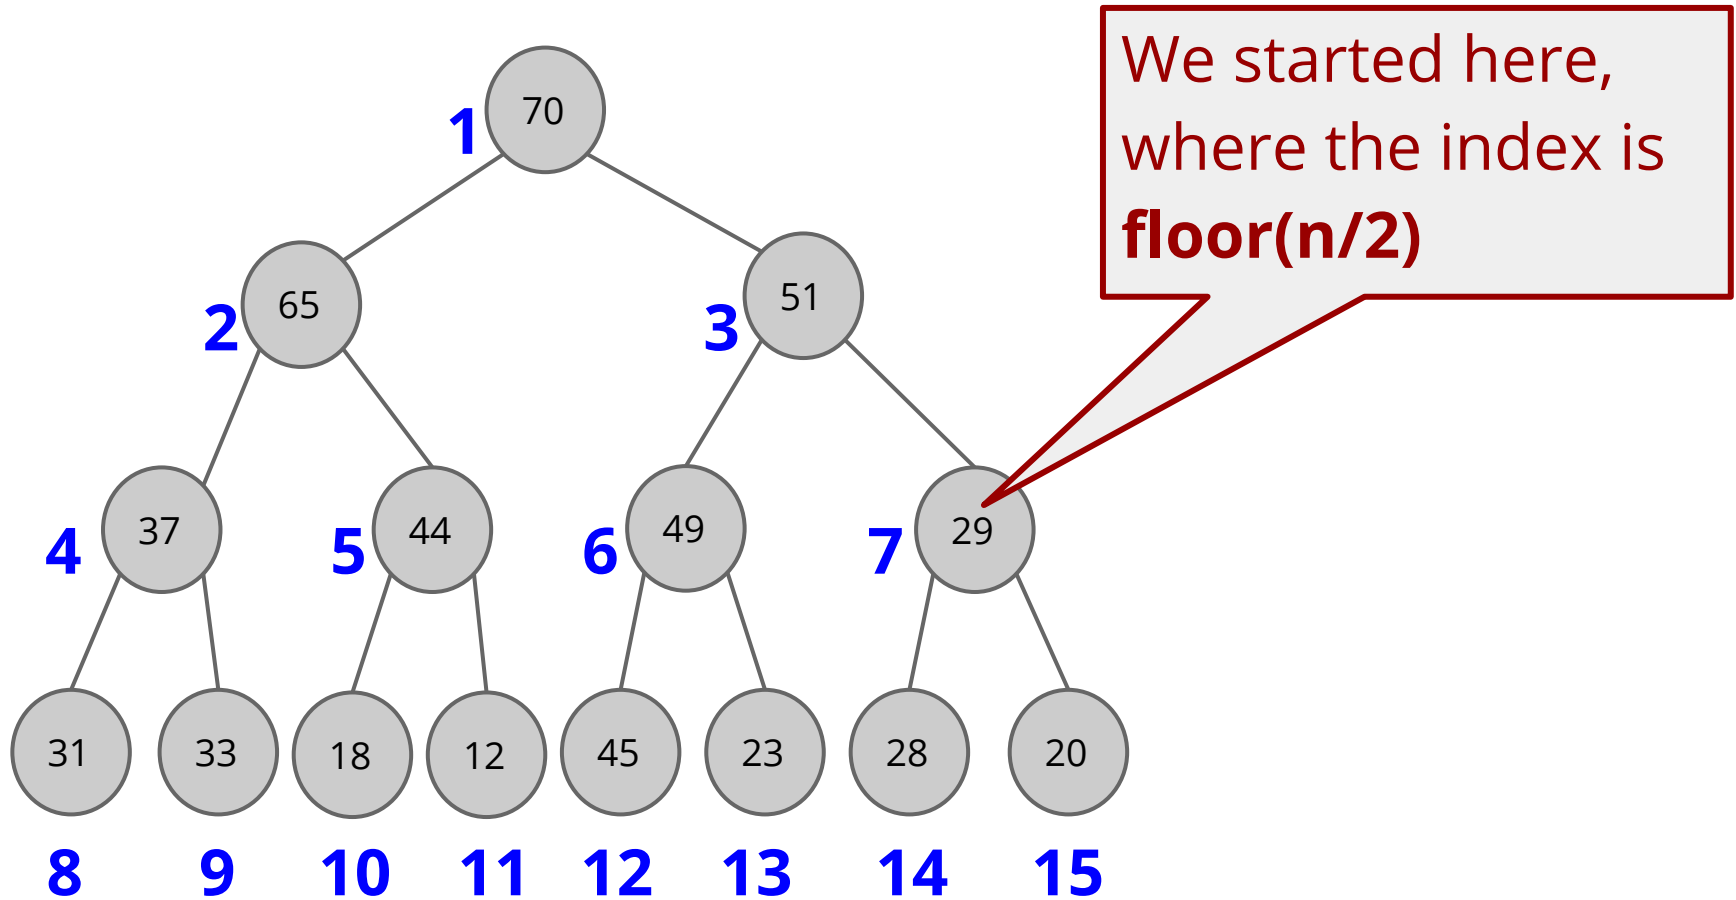
We did nothing but
bubbling-down




Heap Built!



Idea #2: The starting index



Idea #2: The starting index




Even the bottom level is not fully filled, we still start from **floor(n/2)**

We **always** start from **floor(n/2)**, and go down to 1.

Idea #2: Pseudo-code!

```
BuildMaxHeap(A):
```

```
    for  $i \leftarrow \text{floor}(n/2)$  downto 1:  
        BubbleDown(A, i)
```



Advantages of Idea #2:

- It's in-place, no need for extra array (we did nothing but bubble-down, which is basically swappings).
- It's worst-case running time is **$O(n)$** , instead of $O(n \log n)$ of Idea #1.

Why?

Analysis:

Worst-case running time of

BuildMaxHeap(A)



Intuition

A complete binary tree with n nodes...

How many levels?
 $\sim \log n$


$n/16$ nodes, and # of swaps per bubble-down: ≤ 3

$n/8$ nodes, and # of swaps per bubble-down: ≤ 2

of swaps per bubble-down: ≤ 1

$\sim n/4$ nodes

$\sim n/2$ nodes, and **no** work done at this level.



So, total number of swaps

$$T(n) = 1 \cdot \frac{n}{4} + 2 \cdot \frac{n}{8} + 3 \cdot \frac{n}{16} + \dots$$


$$= \sum_{i=1}^{\log n} i \cdot \frac{n}{2^{i+1}} \leq \sum_{i=1}^{+\infty} i \cdot \frac{n}{2^{i+1}}$$

$$= n \sum_{i=1}^{+\infty} \frac{i}{2^{i+1}}$$



same trick as
Week 1's sum

$$= n$$



A close-up of Morpheus from The Matrix, wearing his signature sunglasses. The image is used as a background for a meme. The text is overlaid in a bold, white, sans-serif font with a black outline. The top text is at the top of the image, and the bottom text is at the bottom. The background is a blurred outdoor scene.

BUILD MAX HEAP

YOU CAN DO IN LINEAR TIME

Summary

HeapSort(A):

- Sort an unsorted array in-place
- $O(n \log n)$ worst-case running time

BuildMaxHeap(A):

- Convert an unsorted array into a heap, in-place
- Fix heap property from bottom up, do bubbling down on each sub-root
- $O(n)$ worst-case running time

Algorithm visualizer

<http://visualgo.net/heap.html>

Next week

→ ADT: Dictionary

→ Data structure: Binary Search Tree