

CSC263 Week 9

Larry Zhang

Announcements

- Midterm, class average 62.5% (37.5/60)
- PS7 out soon, due next Tuesday
- A2 out, due March 31, start early!
- Don't forget to give feedback (especially about the midterm)
 - ◆ <http://goo.gl/forms/S9yie3597B>

Recap

→ The Graph ADT

- ◆ definition and data structures

→ BFS

- ◆ gives us single-source **shortest** path
- ◆ Let $\delta(\mathbf{s}, \mathbf{v})$ denote the length of shortest path from \mathbf{s} to \mathbf{v} ...
- ◆ then after performing a BFS starting from \mathbf{s} , we have, for all vertices \mathbf{v}

$$\mathbf{d}[\mathbf{v}] = \delta(\mathbf{s}, \mathbf{v})$$

We can totally prove it.

Idea of the proof

There is no way $d[v] < \delta(s, v)$, according to Lemma 22.2

Use contradiction: suppose there exist v s.t. $d[v] > \delta(s, v)$, let v be the one with the **minimum** $\delta(s, v)$.

Then on a shortest path between s and v , pick vertex u which is immediately before v ...

then we have **$d[v] > \delta(s, v) = \delta(s, u) + 1 = d[u] + 1$**

Must be equal because u is on the shortest path from s to v .

Must be equal because v is the minimum $\delta(s, v)$ that violates $d[v] > \delta(s, v)$, so u must not be violating.

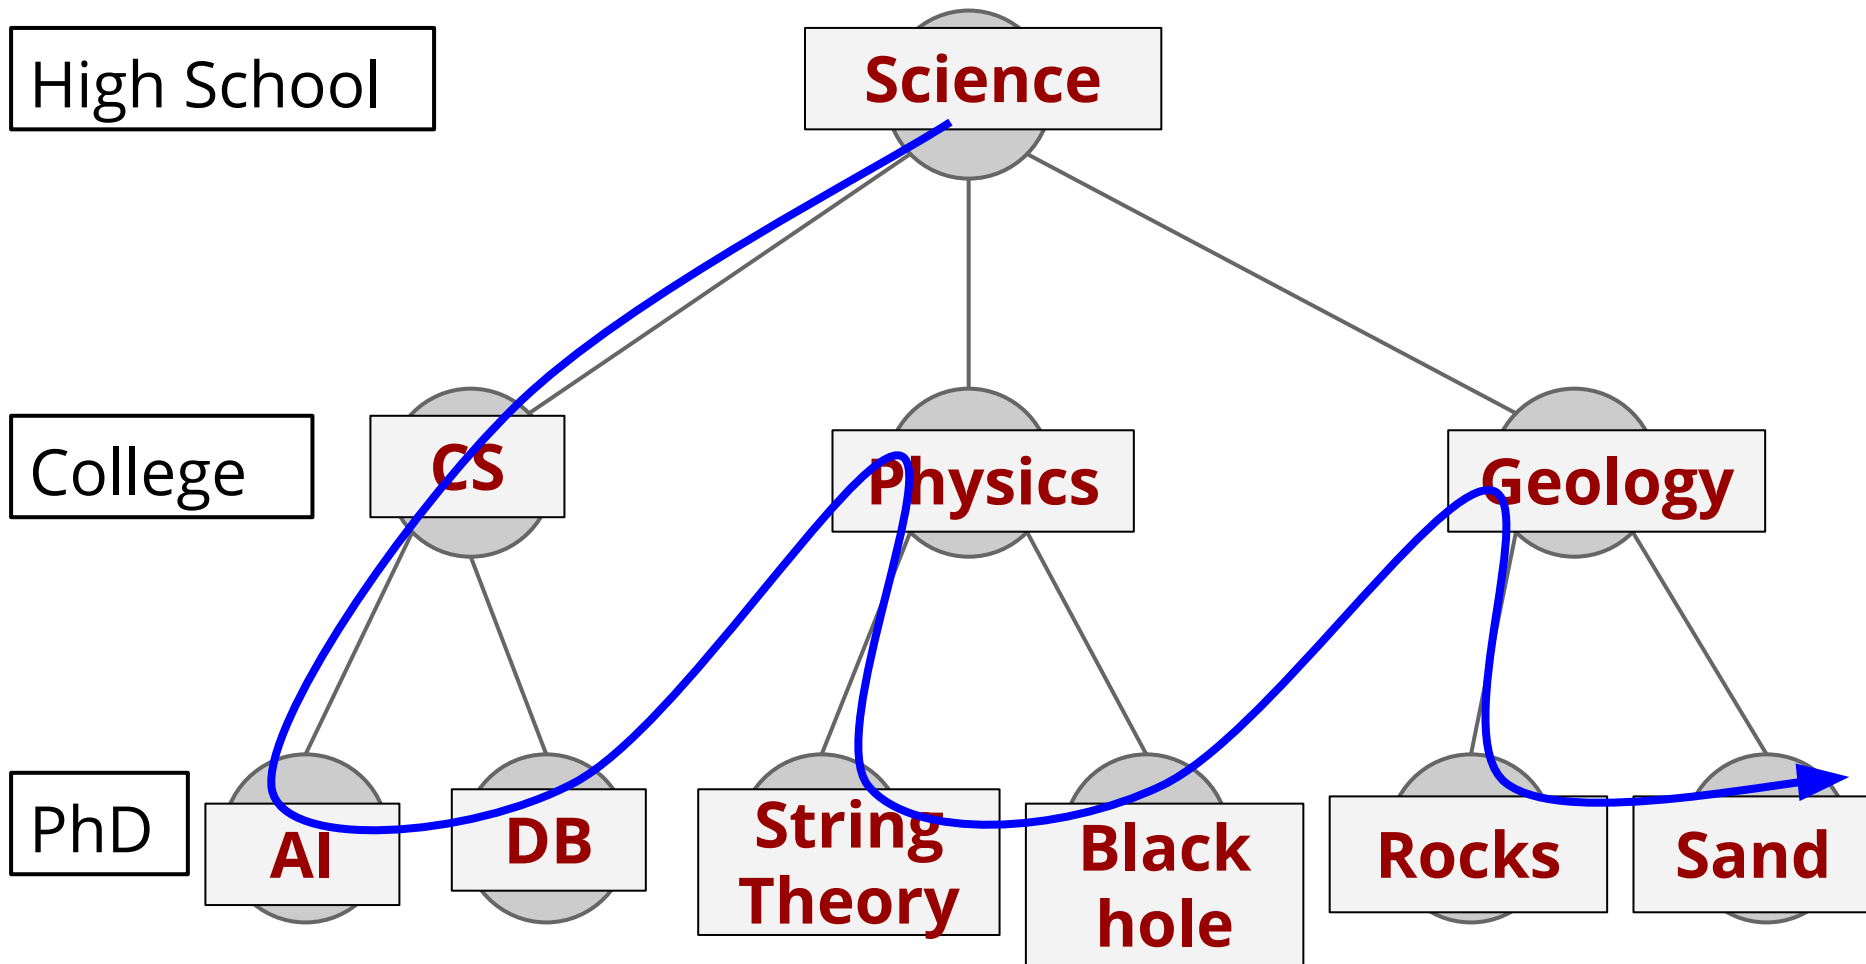
Think about the moment after dequeue u (checking u 's neighbours)

- if v is white, $d[v] = d[u] + 1$ (how BFS works), **contradiction!**
- if v is black, $d[v] \leq d[u]$ (coz v is dequeued before u), **contradiction!**
- if v is gray, then it is coloured gray by some other vertex w , then $d[v] = d[w] + 1$ and $d[w] \leq d[u]$, therefore $d[v] \leq d[u] + 1$, **contradiction!**

Depth-First Search

The **Depth-First** way of learning these subjects

→ Go towards PhD whenever possible; only start learning physics after finishing everything in CS.



DFS



BFS



```
NOT_YET_DFS(root):
```

```
    Q ← Stack()
```

```
    Push(Q, root)
```

```
    while Q not empty:
```

```
        x ← Pop(Q)
```

```
        print x
```

```
        for each child c of x:
```

```
            Push(Q, c)
```

```
NOT_YET_BFS(root):
```

```
    Q ← Queue()
```

```
    Enqueue(Q, root)
```

```
    while Q not empty:
```

```
        x ← Dequeue(Q)
```

```
        print x
```

```
        for each child c of x:
```


```
            Enqueue(Q, c)
```

**Why they are
twins!**

DFS in a tree

Output:

a
c
f
e
b
d



```
NOT_YET_DFS(root):
```

```
  Q ← Stack()
```

```
  Push(Q, root)
```

```
  while Q not empty:
```


```
    x ← Pop(Q)
```

```
    print x
```


```
    for each child c of x:
```

```
      Push(Q, c)
```


Stack:




POP



POP POP



POP POP



POP

A nicer way to write this code?

The use of stack is basically implementing **recursion**.

```
NOT_YET_DFS(root):  
    Q ← Stack()  
    Push(Q, root)  
    while Q not empty:  
        x ← Pop(Q)  
        print x  
        for each child c of x:  
            Push(Q, c)
```

```
NOT_YET_DFS(root):  
    print root  
    for each child c of x:  
        NOT_YET_DFS(c)
```

Exercise: Try this code on the tree in the previous slide.

Avoid visiting a vertex twice, **same as BFS**

Remember you visited it by **labelling** it using **colours**.

- **White**: “unvisited”
- **Gray**: “encountered”
- **Black**: “explored”



- Initially all vertices are **white**
- Colour a vertex **gray** the **first** time visiting it
- Colour a vertex **black** when **all** its **neighbours** have been encountered
- Avoid visiting **gray** or **black** vertices
- In the end, all vertices are **black**

Other values to remember, some are same as BFS

- **pi[v]**: the vertex from which v is encountered
 - ◆ “I was introduced as **whose** neighbour?”

Other values to remember, different from BFS

- There is a **clock** ticking, incremented whenever someone's colour is changed
- For each vertex v , remember two **timestamps**
 - ◆ **$d[v]$** : “discovery time”, when the vertex is first encountered
 - ◆ **$f[v]$** : “finishing time”, when all the vertex's neighbours have been visited.

Note : this $d[v]$ is totally different from that distance value $d[v]$ in BFS!

The pseudo-code (incomplete)


```
DFS_VISIT(G, u):  
    colour[u] ← gray  
    time ← time + 1  
    d[u] ← time      # keep discovery time  
                     # on first encounter  
    for each neighbour v of u:  
        if colour[v] = white:  
            pi[v] ← u  
            DFS_VISIT(G, v)  
    colour[u] ← black  
    time ← time + 1  
    f[u] ← time      # keep finishing time after  
                     # exploring all neighbours
```

The red part is
the same as
NOT_YET_DFS


Why **DFS_VISIT**
instead of **DFS**?
We will see...

Let's run an example!


DFS_VISIT(G, u)




time = 1, encounter the source vertex




time = 2, recursive call, level 2




time = 3, recursive call, level 3




time = 4, recursive call, level 4




time = 5, vertex x finished




time = 6, recursion back to level 3, finish y




time = 7, recursive back to level 2, finish v



time = 8, recursion back to level 1, finish u



DFS_VISIT(G, u) done!



What about these two white vertices?

We actually want to visit them (for some reason)

The pseudo-code for visiting everyone

DFS(G):

for each v in $G.V$:

$\text{colour}[v] \leftarrow \text{white}$

$f[v] \leftarrow d[v] \leftarrow \infty$

$\text{pi}[v] \leftarrow \text{NIL}$

time $\leftarrow 0$

for each v in $G.V$:

 if $\text{colour}[v] = \text{white}$:

 DFS_VISIT(G, v)

Initialization

DFS_VISIT(G, u):

$\text{colour}[u] \leftarrow \text{gray}$

 time $\leftarrow \text{time} + 1$

$d[u] \leftarrow \text{time}$

 for each neighbour v of u :

 if $\text{colour}[v] = \text{white}$:

$\text{pi}[v] \leftarrow u$

 DFS_VISIT(G, v)


$\text{colour}[u] \leftarrow \text{black}$

 time $\leftarrow \text{time} + 1$


$f[u] \leftarrow \text{time}$

Make sure NO vertex is left with white colour.


So, let's finish this DFS




time = 9, DFS_VISIT(G, w)




time = 10




time = 11



time = 12



DFS(G) done!



Runtime analysis!

The total amount of work (use **adjacency list**):

- Visit each vertex once
 - ◆ constant work per vertex
 - ◆ in total: $O(|V|)$
- At each vertex, check all its neighbours (all its **incident edges**)
 - ◆ Each edge is checked **once** (in a directed graph)
 - ◆ in total: $O(|E|)$

Same as BFS

Total runtime:
 $O(|V| + |E|)$

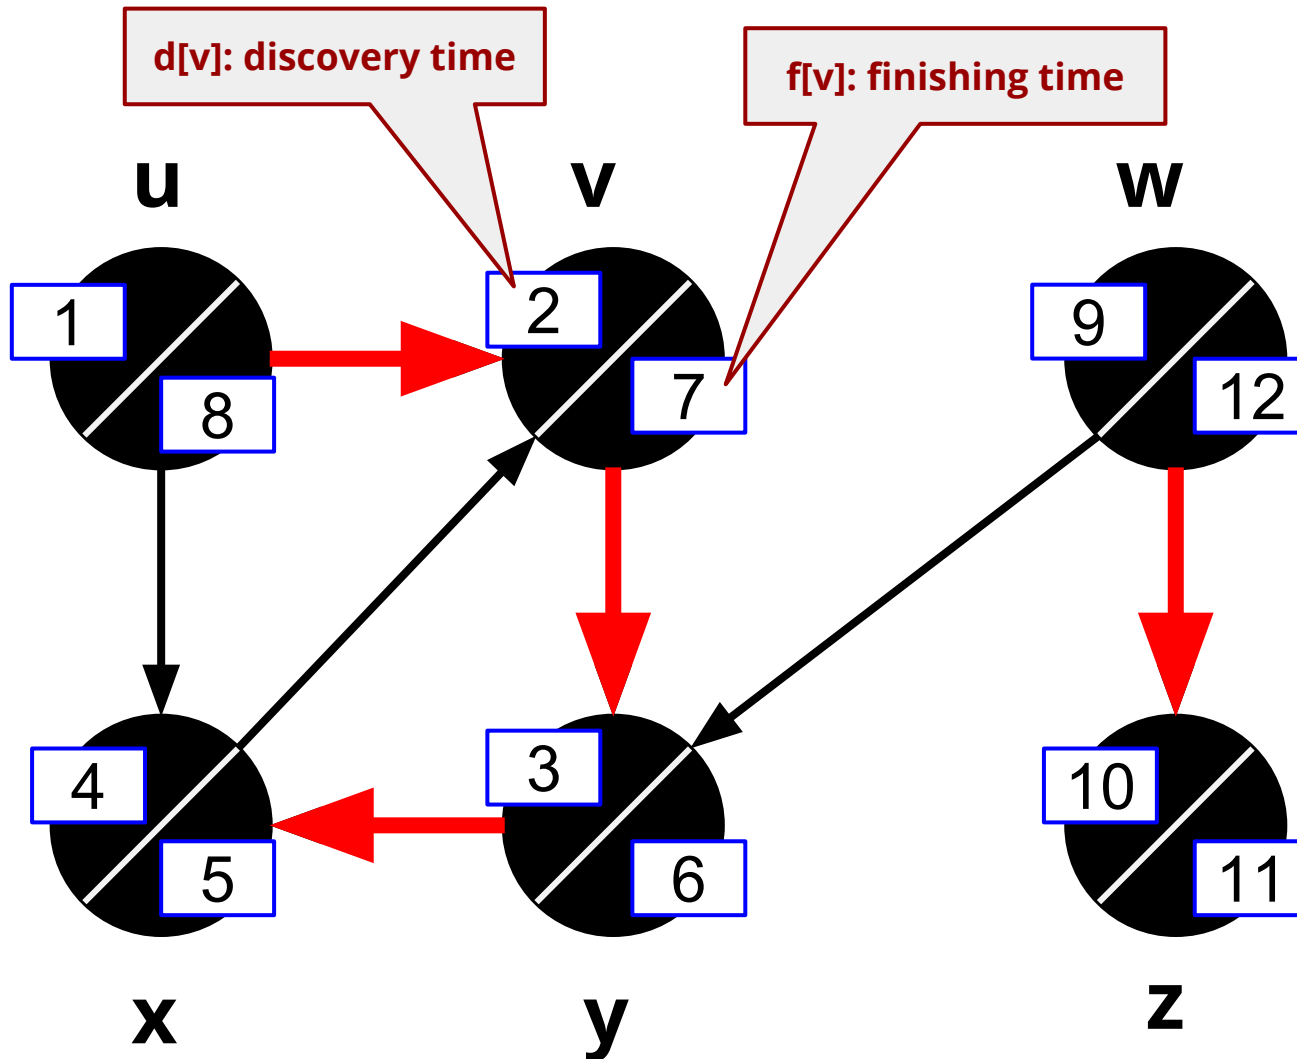
What do we get from DFS?

- Detect whether a graph has a cycle.
 - ◆ That's why we wanted to visit all vertices -- if you want to be sure whether a graph has a cycle or not, you'd better check **everywhere**.
 - ◆ Why didn't we do the similar thing for BFS?
- How exactly do we detect a cycle?

CSC263 Week 9

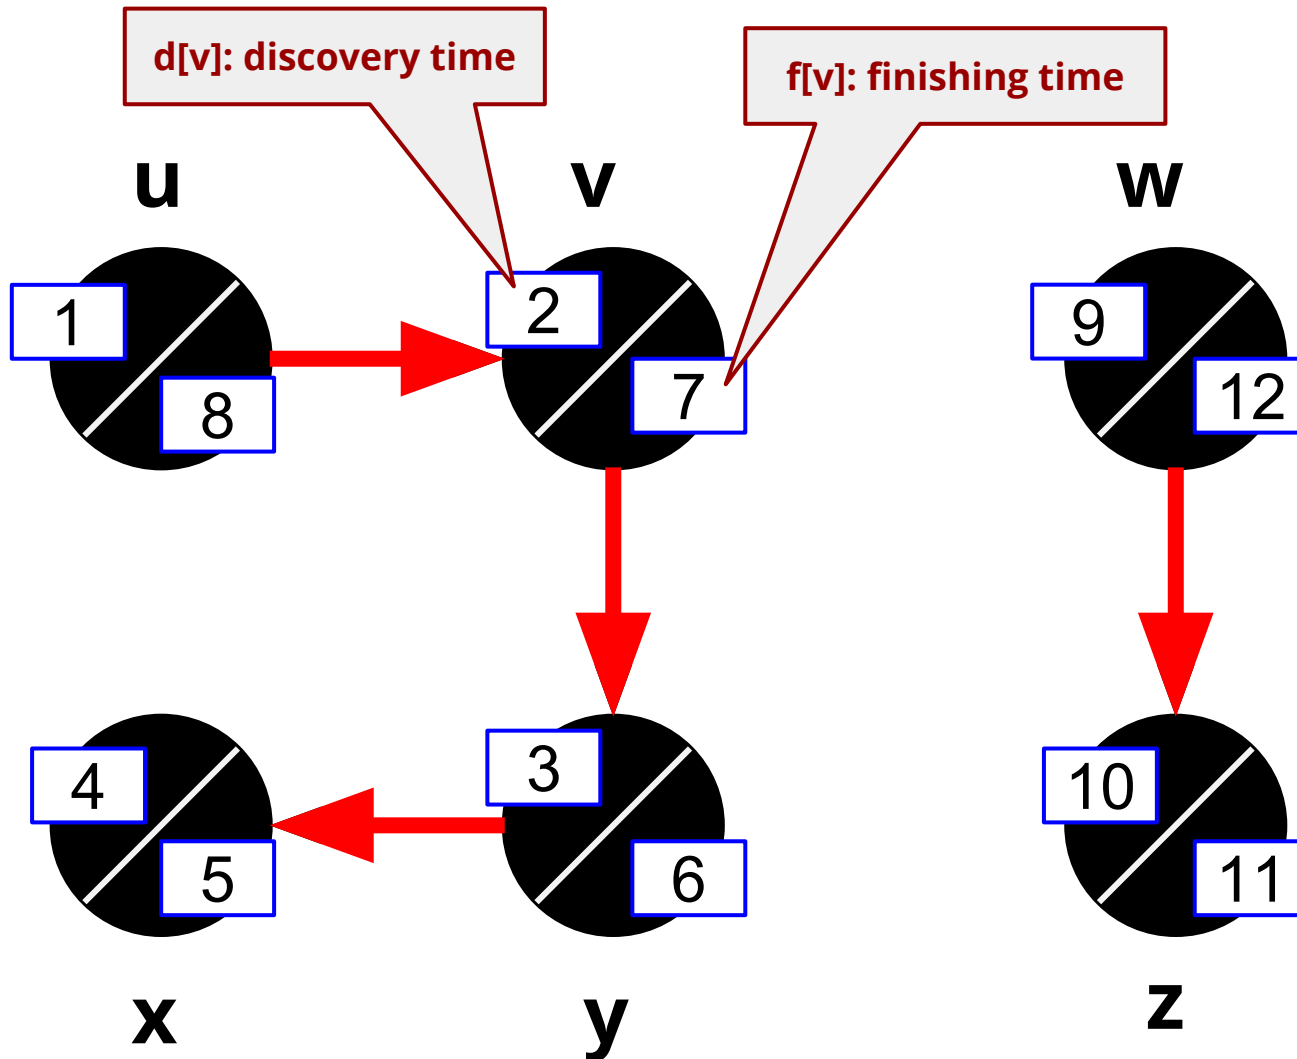
Thursday

Recap: DFS(G) done!



How do we use all the info?


We get a **DFS forest**
(a set of disjoint trees)



**determine descendant / ancestor
relationship in the DFS forest**

How to decide whether **y** is a **descendant** of **u** in the DFS forest?

Idea #1: trace back the **pi[v]** pointers (the red edges) starting from **y**, see whether you can get to **u**.
Worst-case takes **O(n)** steps.



the “parenthesis structure”


((())) () (())

- Either one pair **contains** the another pair.
- Or one pair is **disjoint** from another

(())

This (overlapping)
never happens!



Visualize $d[v]$, $f[v]$ as interval $[d[v], f[v]]$



Now, visualize all the intervals!

What do you see in this?

Parenthesis structure!



The $[d[v], f[v]]$ intervals that we got from DFS follow the parenthesis structure, i.e.,

- Either one interval **contains** another
- Or one is **disjoint** from another


Moreover,


- **Iff** interval of **u** contains interval of **v**, then **u** is an **ancestor** of **v** in the DFS forest.
- If interval of **u** is disjoint from interval of **v**, then they are **not** ancestors of each other.

How to decide whether **y** is a **descendant** of **u** in the DFS forest?

Idea #1: trace back from **y** to **u**.
(the red arrows) takes $O(n)$ steps.

Idea #2: see if $[d[u], f[u]]$ contains $[d[y], f[y]]$.
Worst-case: **1 step!**





We can efficiently check whether a vertex is an ancestor of another vertex in the DFS forest.


so what...



Classifying Edges

4 types of edges in a graph after a DFS


- **Tree edge:** an edge in the DFS-forest
- **Back edge:** a non-tree edge pointing from a vertex to its **ancestor** in the DFS forest.
- **Forward edge:** a non-tree edge pointing from a vertex to its **descendant** in the DFS forest




→ **Cross edge:** all other edges

Checking edge types


We can efficiently check edge types, because...
we can efficiently check whether a vertex is an **ancestor / descendant** of another vertex using...
the **parenthesis structure** of $[d[v], f[v]]$ intervals!






We can efficiently check edge types after a DFS!


so what...




A graph is **cyclic** if and only if DFS yields a **back edge**.



That's useful!



A (directed) graph contains a **cycle** if and only if DFS yields a **back edge**




A (directed) graph contains a **cycle** if and only if DFS yields a **back edge**

Proof of **"if"**:

Let the edge be (u, v) ,
then by definition of back edge, v is an ancestor of u in the DFS tree,
then there is a path from v to u , i.e., $v \rightarrow \dots \rightarrow u$,
plus the back edge $u \rightarrow v$,
BOOM! Cycle.

Proof of **"only if"**:

Let the cycle be...,




Let v_0 be the first one that turns gray, when all others in the cycle are white, then v_k must be a descendant of v_0 .
(Read "White Path Theorem" in Text)

How about undirected graph?


Should **back** and **forward** edges be the same thing?

→ No, because although the edges are undirected, **neighbour checking** still has a “direction”.




More about undirected graph

After a DFS on a undirected graph, **every** edge is either a **tree edge** or a **back edge**, i.e., **no** forward edge or cross edge.



If this were a forward edge, it would violate the DFS algorithm (not checking at C but tracing back and check at A)

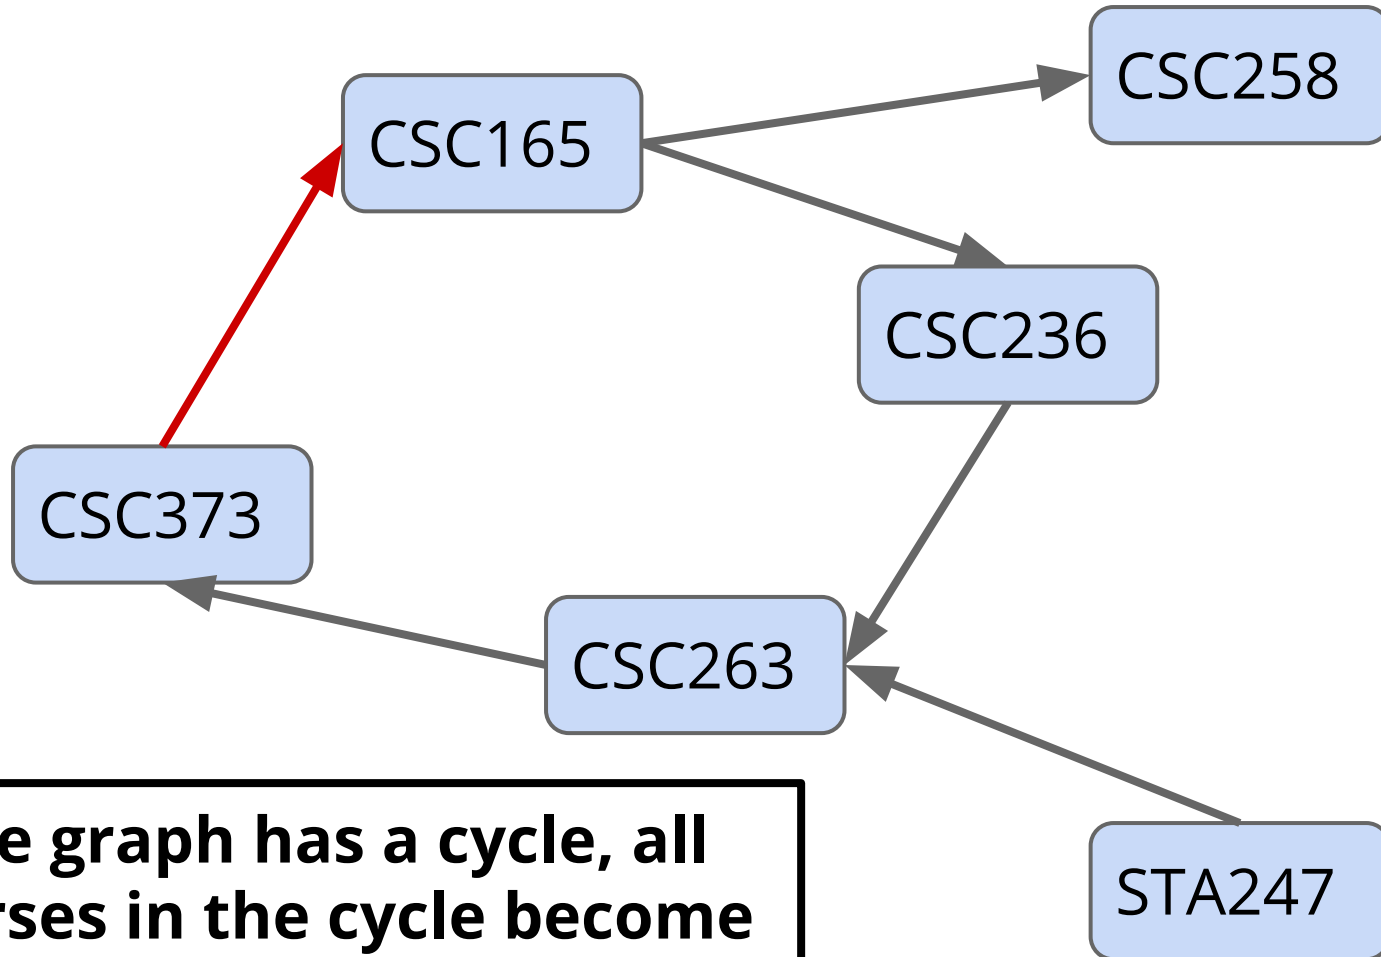


If this were a cross edge, it violates DFS again (should have checked (A, C) when reached A, but instead wait until C is visited.)

Why do we care about **cycles** in a graph?

Because cycles have meaningful
implication in real applications.

Example: a course prerequisite graph




If the graph has a cycle, all
courses in the cycle become
impossible to take!

Applications of DFS


- Detect cycles in a graph
- Topological sort
- Strongly connected components

Topological Sort

→ Place the vertices in such an order that all edges are pointing to the right side.



A valid order of getting dressed.




How to do topological sorting

1. Do a **DFS**
2. Order vertices according to their **finishing times $f[v]$**

Strongly connected components

- Subgraphs with strong connectivity (any pair of vertex can reach each other)



Summary of DFS

- It's the twin of BFS (Queue vs Stack)
- Keeps two timestamps: $d[v]$ and $f[v]$
- Has same runtime as BFS
- Does NOT give us shortest-path
- Give us cycle detection (back edge)
- For real problems, choose BFS and DFS wisely.

Next week

→ Minimum Spanning Tree



<http://goo.gl/forms/S9yie3597B>