

CSC263 Week 11

Larry Zhang

<http://goo.gl/forms/S9yie3597B>

Announcements

→ A2 due next Tuesday

→ Course evaluation:

<http://uoft.me/course-evals>

ADT: Disjoint Sets


- What does it store?
- What operations are supported?

What does it store?

The elements in the sets can change dynamically.


It stores a collection of (**dynamic**) **sets** of elements, which are **disjoint** from each other.

Each element belongs to **only one** set.



Each set has a **representative**


A set is **identified** by its representative.



Operations

MakeSet(x): Given an element x that does NOT belong to any set, create a new set $\{x\}$, that contains only x , and assign x as the representative.

MakeSet("Newton")




Operations

FindSet(x): return the representative of the set that contains **x**.

FindSet("Bieber") returns: **Ford**

FindSet("Oprah") returns: **Obama**

FindSet("Newton")
returns: **Newton**





Operations

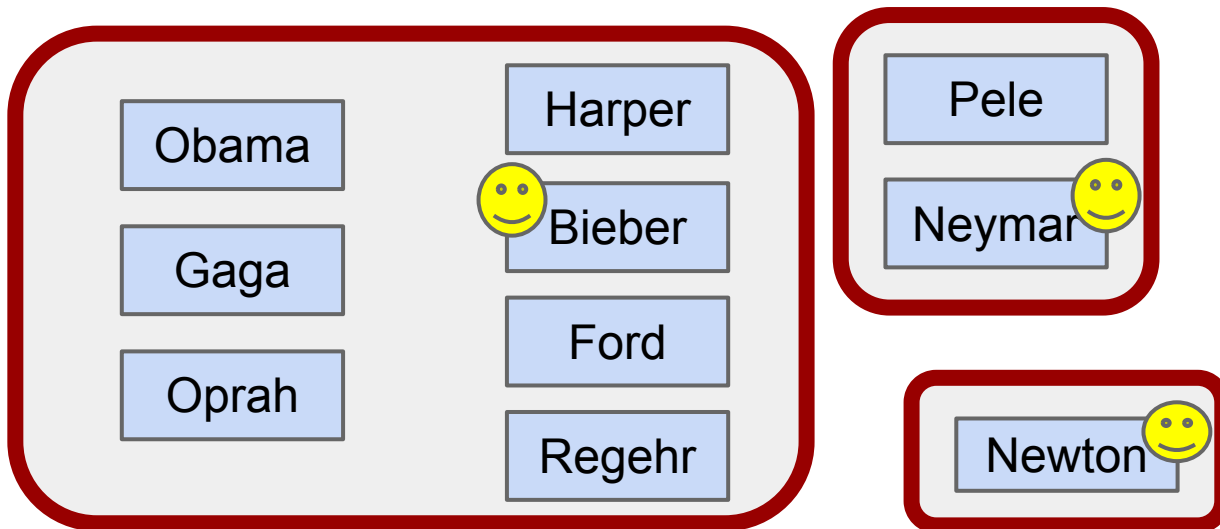
If **x** and **y** are already in the **same** set, then nothing happens.

Union(x, y): given two elements **x** and **y**, create a **new set** which is the **union** of the two sets that contain **x** and **y**, **delete** the original sets that contains x and y.

Pick a **representative** of the new set, usually (but not necessarily) one of the representatives of the two original sets.



Union("Gaga", "Harper")



Applications


KRUSKAL-MST($G(V, E, w)$):

```
1  T ← {}
2  sort edges so that  $w(e_1) \leq w(e_2) \leq \dots \leq w(e_m)$ 
3  for each v in V:
4      MakeSet(v)
5  for i ← 1 to m:
6      # let  $(u_i, v_i) = e_i$ 
7      if FindSet( $u_i$ ) != FindSet( $v_i$ ):
8          Union( $u_i, v_i$ )
9      T ← T ∪ { $e_i$ }
```

Other applications

For each edge (u, v)
if $\text{FindSet}(u) \neq \text{FindSet}(v)$,
then $\text{Union}(u, v)$

Finding connected components of a graph



Summary: the ADT

- Stores a collection of disjoint sets
- Supported operations
 - ◆ MakeSet(x)
 - ◆ FindSet(x)
 - ◆ Union(x, y)


How to **implement the
Disjoint Sets ADT (efficiently) ?**

Ways of implementations

1. Circularly-linked lists
2. Linked lists with extra pointer
3. Linked lists with extra pointer and with union-by-weight
4. Trees
5. Trees with union-by-rank
6. Trees with path-compression
7. Trees with union-by-weight and path-compression


Circularly-linked list

Circularly-linked list



- One circularly-linked list per set
- Head of the linked list also serves as the representative.

Circularly-linked list



→ **MakeSet(x)**: just a new linked list with a single element x

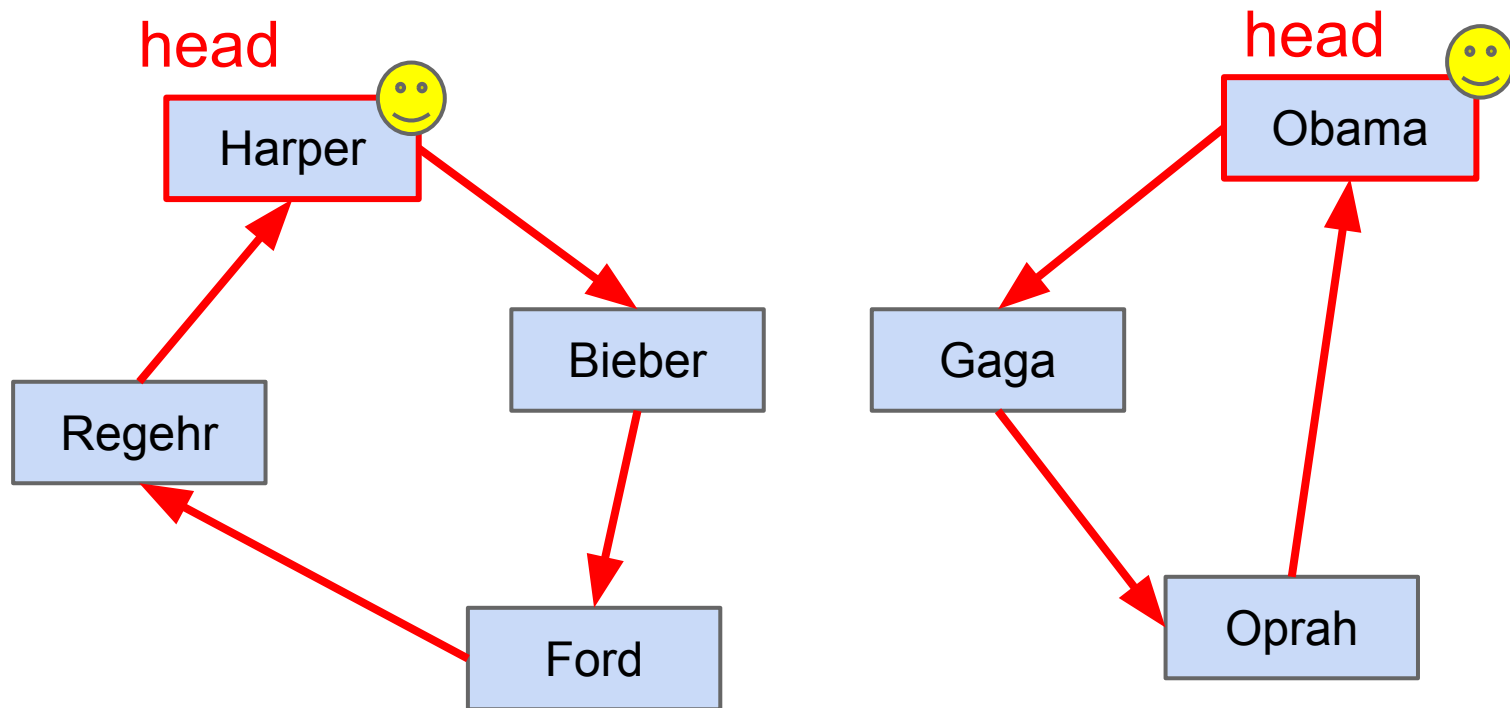
◆ worst-case: **$O(1)$**

→ **FindSet(x)**: follow the links until reaching the head

◆ **$\Theta(\text{Length of list})$**

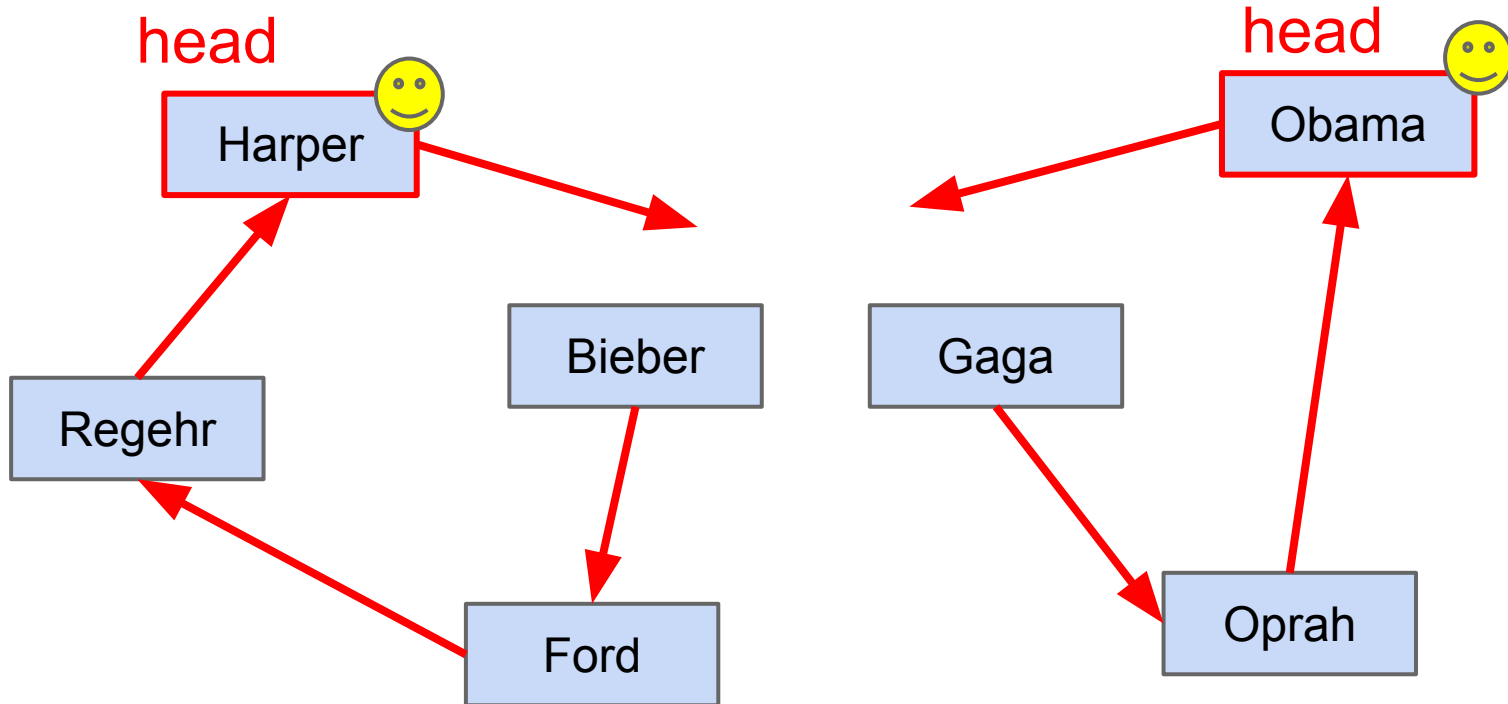
→ **Union(x, y): ...**

Circularly-linked list: **Union(Bieber, Gaga)**

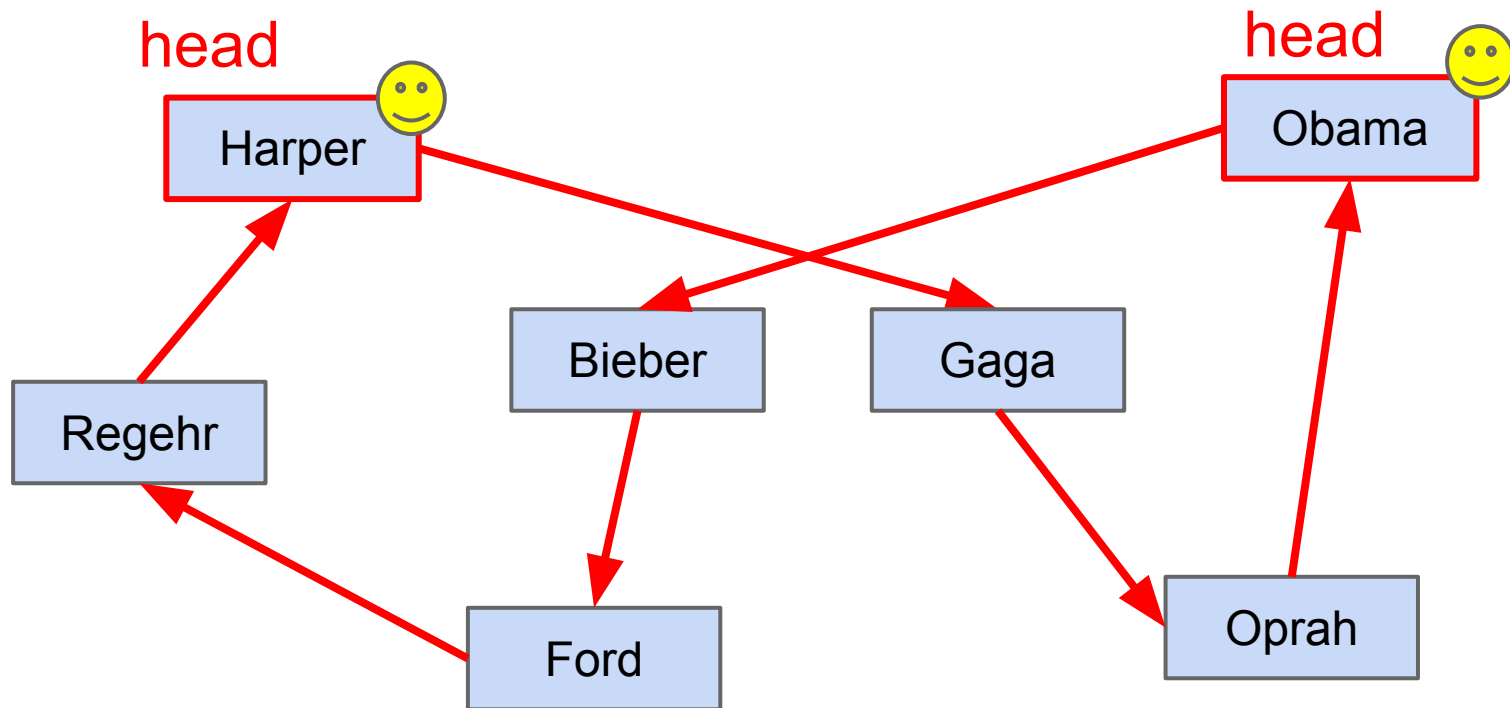


First, locate the head of each linked-list by calling FindSet, takes $\Theta(L)$

Circularly-linked list: **Union...** 1

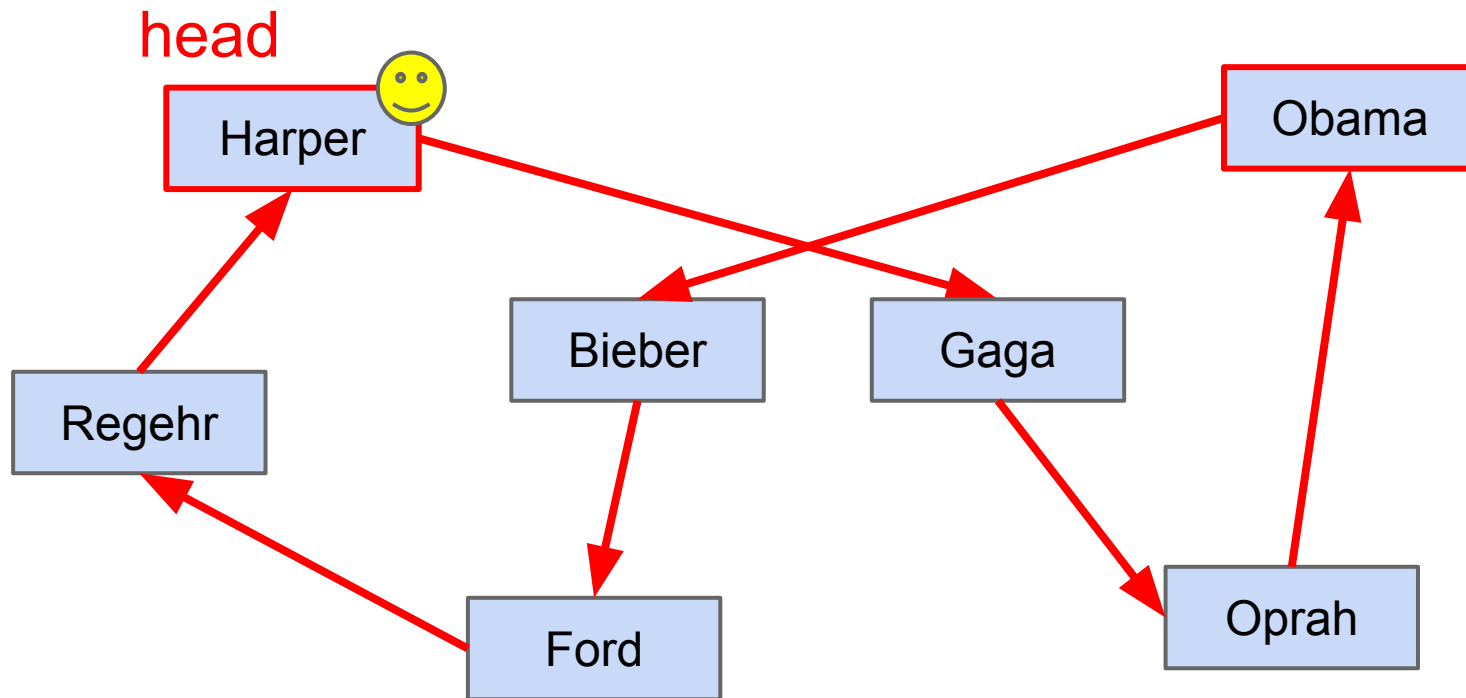


Circularly-linked list: Union... 2



Exchange the two heads' "next" pointers, **$O(1)$**

Circularly-linked list: **Union... 3**



Keep only one representative for the new set.

Circularly-linked list: runtime


FindSet is the time consuming operation

Amortized analysis: How about the **total cost** of a sequence of **m** operations (MakeSet, FindSet, Union)?

- A bad sequence: **$m/4$** MakeSet, then **$m/4 - 1$** Union, then **$m/2 + 1$** FindSet
 - ◆ why it's bad: because many FindSet on a large set (of size $m/4$)
- Total cost: **$\Theta(m^2)$**
 - ◆ each of the **$m/2 + 1$** FindSet takes **$\Theta(m/4)$**

**Linked list
with extra pointer
(to head)**


Linked list with **pointer to head**



- **MakeSet** takes **$O(1)$**
- **FindSet** now takes **$O(1)$** , since we can go to head in 1 step, better than circular linked list
- **Union...**


Linked list with **pointer to head**

Union(Bieber, Pele)



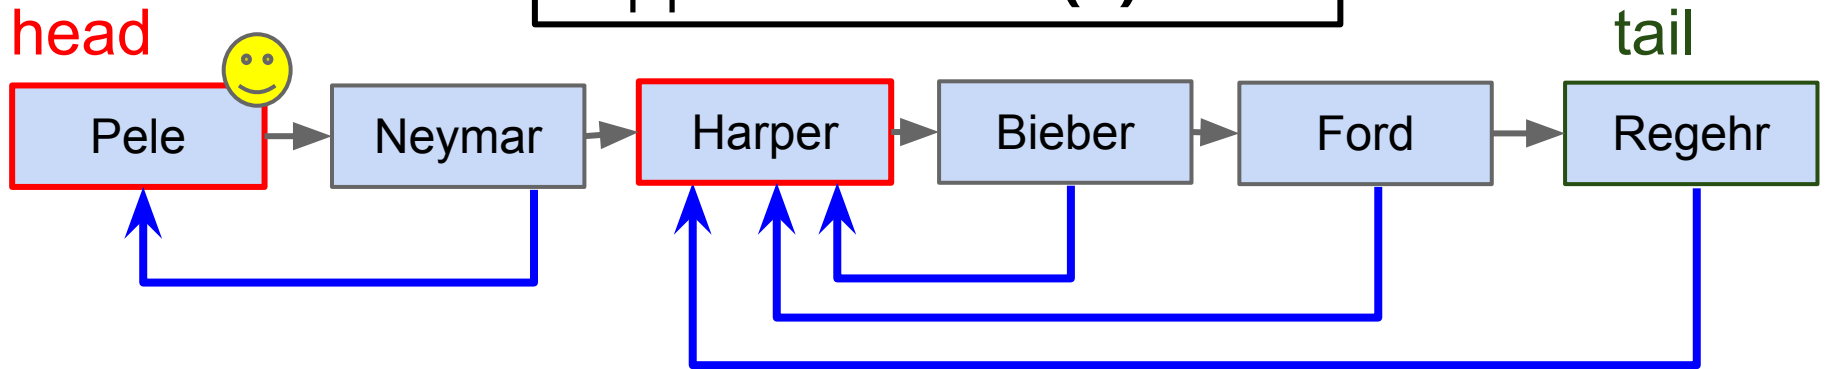
Idea:

Append one list to the other, then **update** the pointers to head

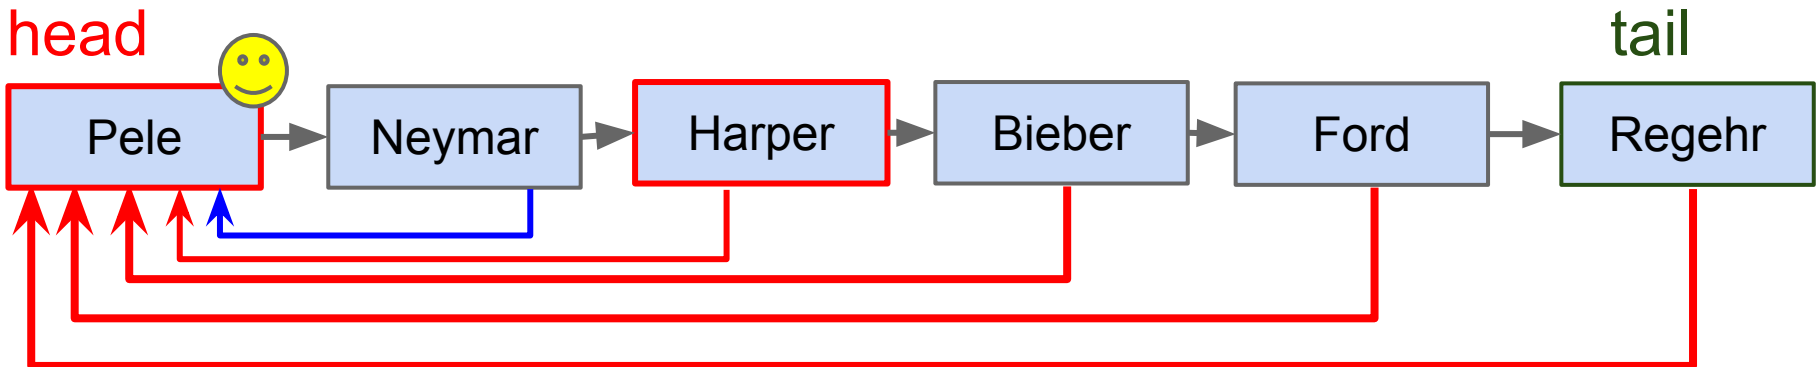


Linked list with **pointer to head**

Append takes **$O(1)$** time



Update **pointers** take **$O(L \text{ of appending list})$**



Linked list with **pointer to head**

MakeSet and **FindSet** are fast, **Union** now becomes the time-consuming one, especially if appending a long list.



Amortized analysis: The total cost of a sequence of **m** operations.

- Bad sequence: **m/2** MakeSet, then **m/2 - 1** Union, then 1 whatever.
 - ◆ Always let the longer list append, like 1 appd 1, 2 appd 1, 3 appd 1, ..., m/2 - 1 appd 1.
- Total cost: $\Theta(1+2+3+\dots+m/2 - 1) = \Theta(m^2)$

Linked list
with extra pointer to head
with union-by-weight


Linked list with **union-by-weight**

Union(Bieber, Pele)

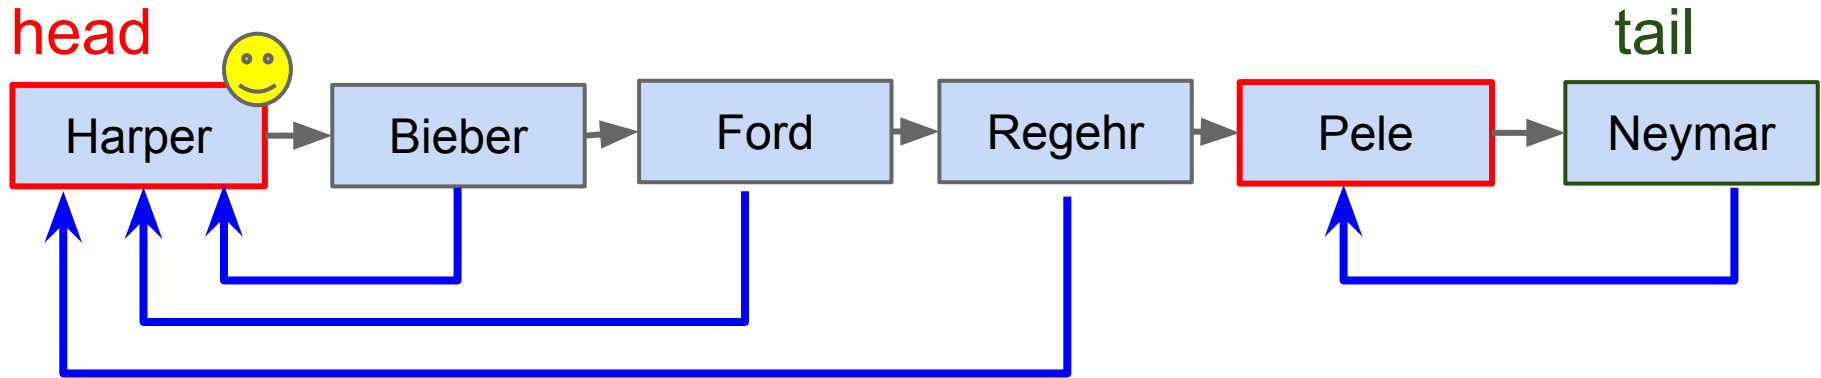


Here we have a choice, let's be a bit smart about it...

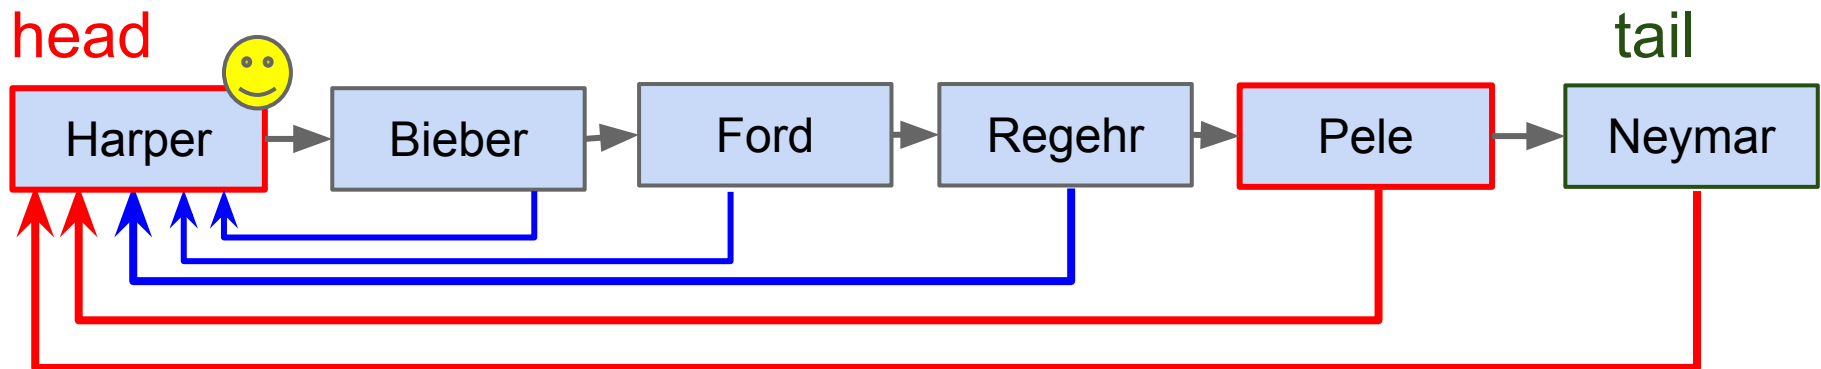
Append the shorter one to the longer one



Linked list with **union-by-weight**



Need to keep track of the **size (weight)** of each list, therefore called **union-by-weight**



Linked list with **union-by-weight**

Union-by-weight sounds like a simple heuristic, but it actually provides significant improvement.

For a sequence of **m** operations which includes **n** MakeSet operations, i.e., **n** elements in total, the total cost is **$O(m + n \log n)$**

i.e., for the previous sequence with $m/2$ MakeSet and $m/2 - 1$ Union, the total cost would be **$O(m \log m)$** , as opposed to **$\Theta(m^2)$** when without union-by-weight.

Linked list with **union-by-weight**

Proof: (assume there are n elements in total)

- Consider an arbitrary element x , how many times does its head pointer need to be updated?
- Because **union-by-weight**, when x is updated, it must be in the smaller list of the two. In other words, after **union**, the size of list at least **doubles**.
- That is, every time x is **updated**, set size **doubles**.
- There are only n elements in total, so we can double at most **$O(\log n)$** times, i.e., x can be updated at most **$O(\log n)$** .
- Same for all n elements, so total updates **$O(n \log n)$**

CSC263 Week 11

Thursday

Ways of implementing Disjoint Sets

- ✓ 1. Circularly-linked lists $\Theta(m^2)$
 - ✓ 2. Linked lists with extra pointer $\Theta(m^2)$
 - ✓ 3. Linked lists with extra pointer and with union-by-weight $\Theta(m \log m)$
- 4. Trees
 - 5. Trees with union-by-rank
 - 6. Trees with path-compression
 - 7. Trees with union-by-weight and path-compression

Benchmark:

Worst-case
total cost of a
sequence of m
operations
(MakeSet or
FindSet or Union)

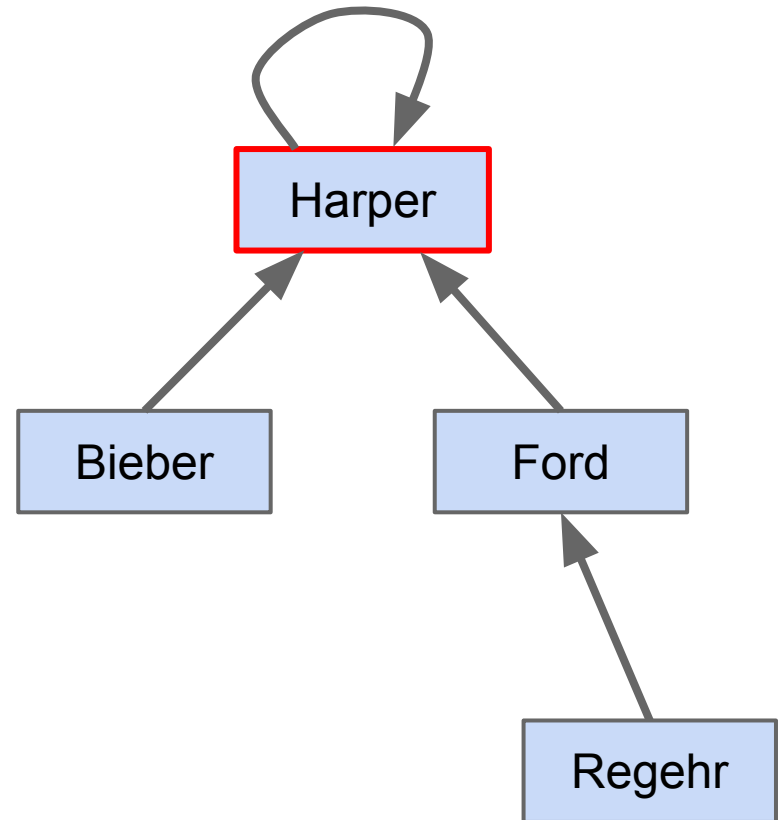
Trees

a.k.a. disjoint set forest




Each set is an “inverted” tree

- Each element keeps a pointer to its **parent** in the tree
- The root points to **itself** (test root by **$x.p = x$**)
- The **representative** is the root
- NOT necessarily a binary tree or balanced tree



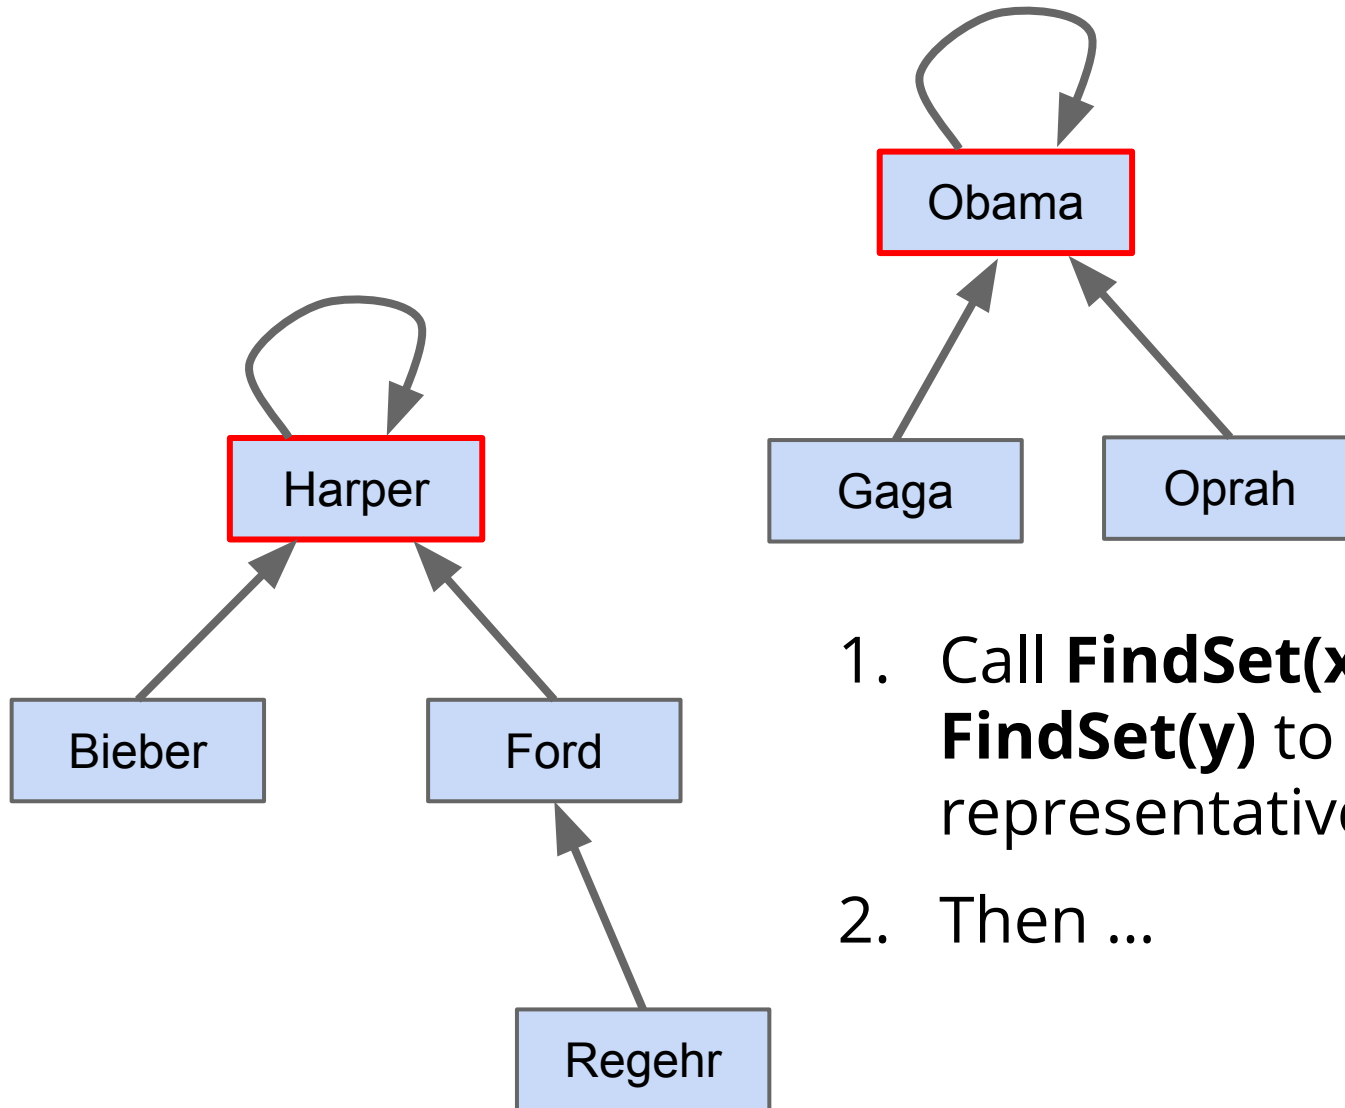
Operations

- **MakeSet(x)**: create a single-node tree with root x
- ◆ **$O(1)$**
- **FindSet(x)**: Trace up the parent pointer until the root is reached
- ◆ **$O(\text{height of tree})$**
- **Union(x, y)...**



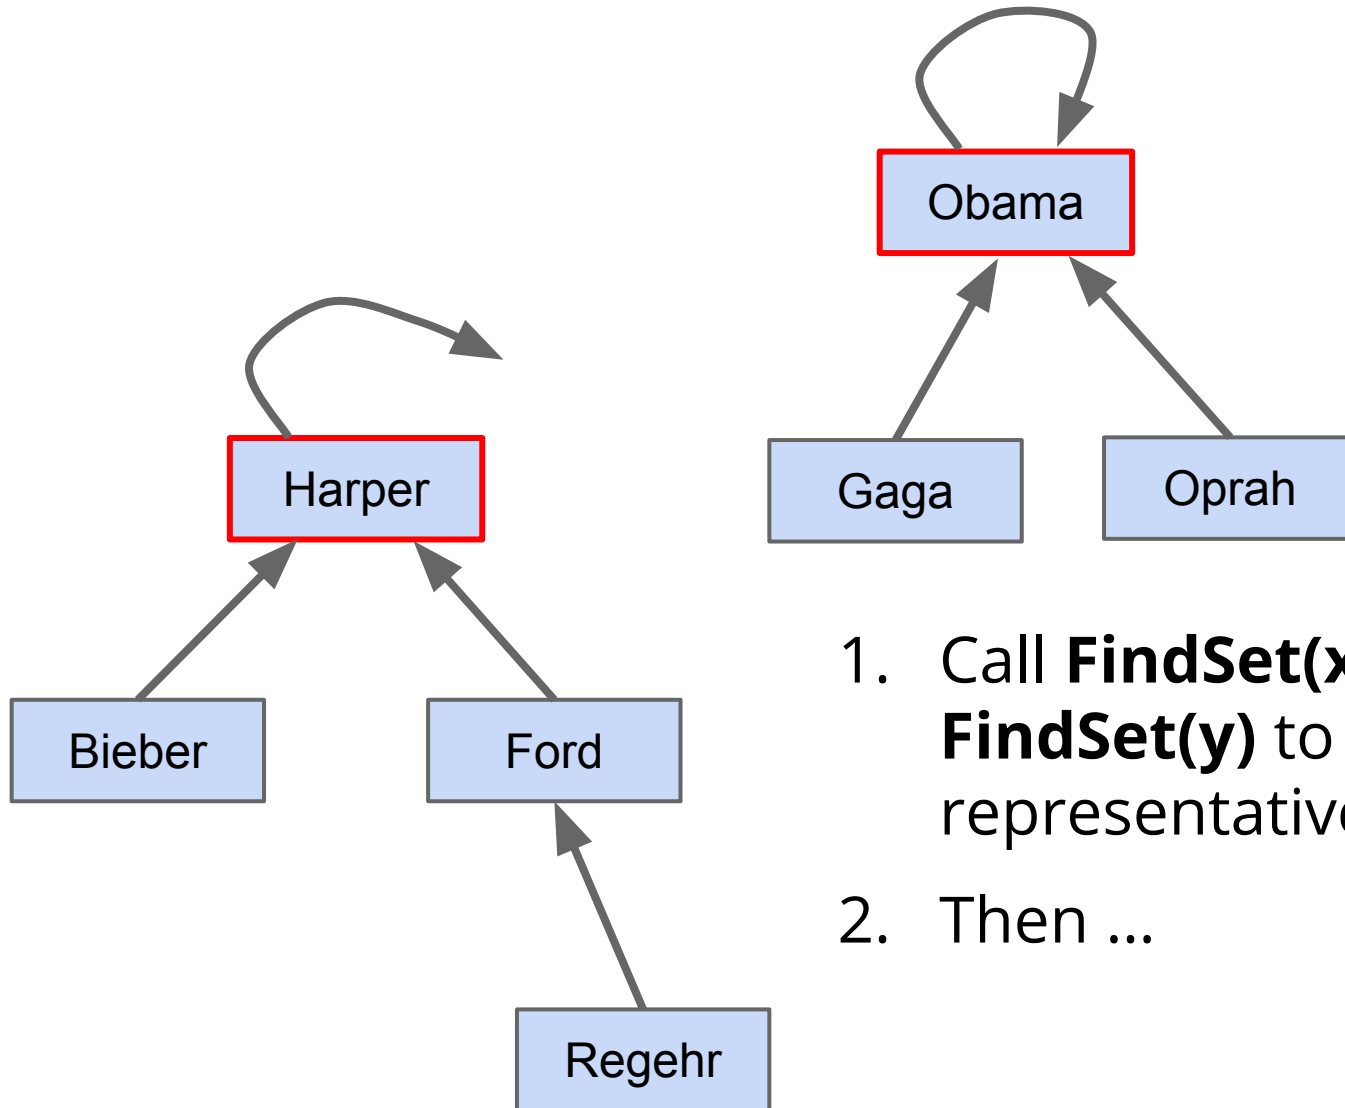
Trees with small heights would be nice.

Union(Bieber, Gaga)



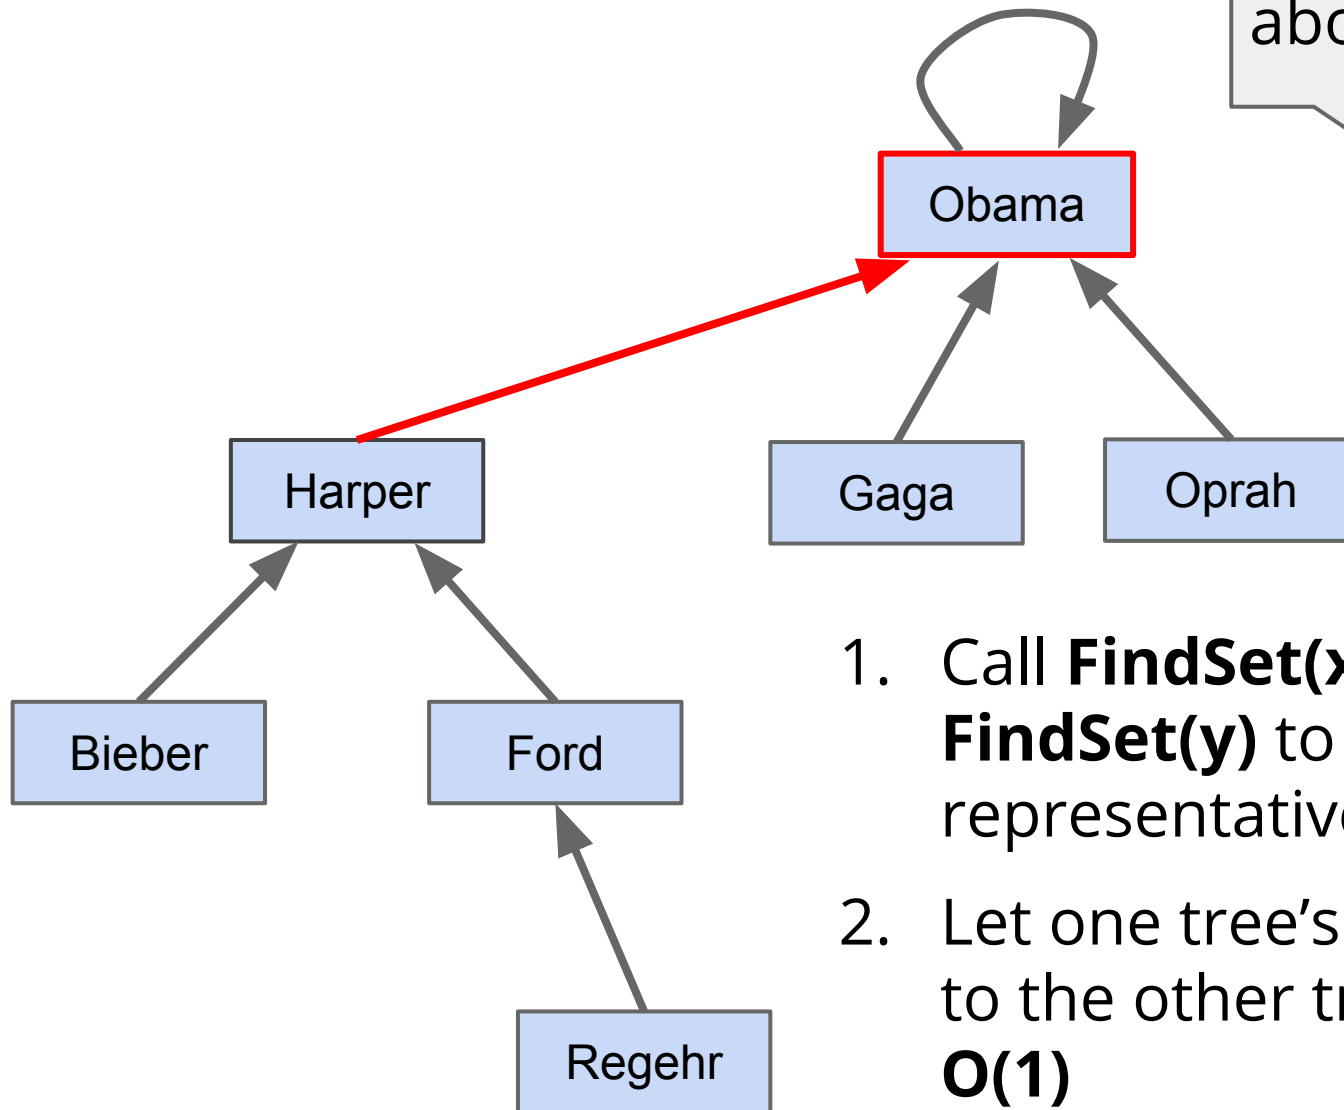
1. Call **FindSet(x)** and **FindSet(y)** to locate the representatives, **$O(h)$**
2. Then ...

Union(Bieber, Gaga)




1. Call **FindSet(x)** and **FindSet(y)** to locate the representatives, **$O(h)$**
2. Then ...

Union(Bieber, Gaga)



Could we have been smarter about this?




1. Call **FindSet(x)** and **FindSet(y)** to locate the representatives, **$O(h)$**
2. Let one tree's root point to the other tree's root, **$O(1)$**

Benchmarking: runtime

The worst-case sequence of **m** operations.
(with **FindSet** being the bottleneck)

m/4 MakeSets, **m/4 - 1** Union, **m/2 + 1** FindSet



Total cost in worst-case
sequence :

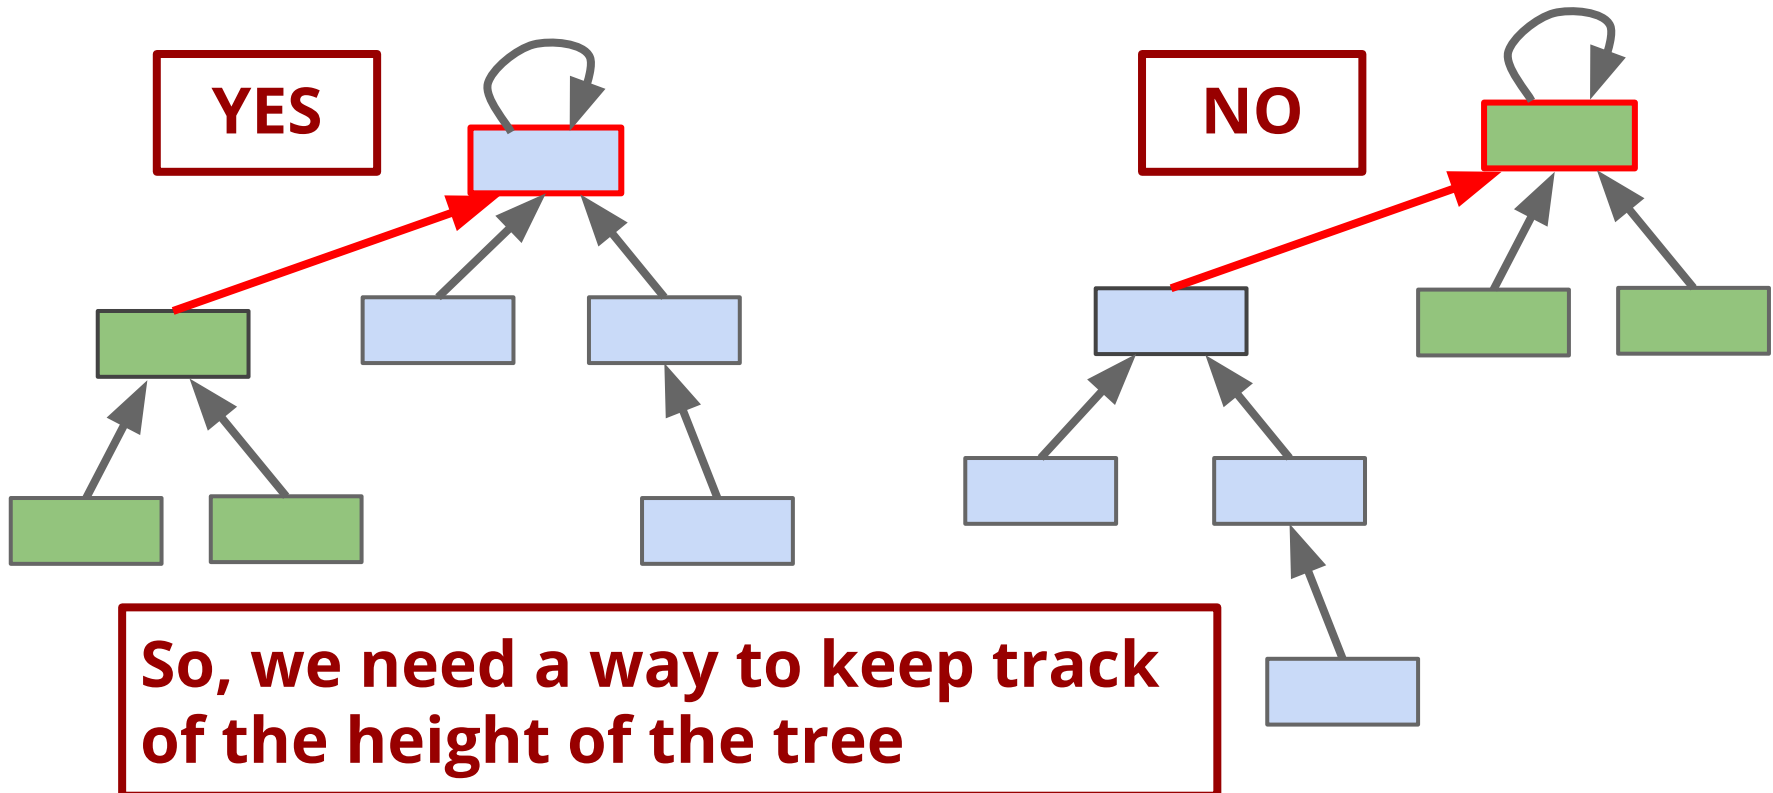
$$\Theta(m^2)$$

(each FindSet would take up to
m/4 steps)

Trees with union-by-rank

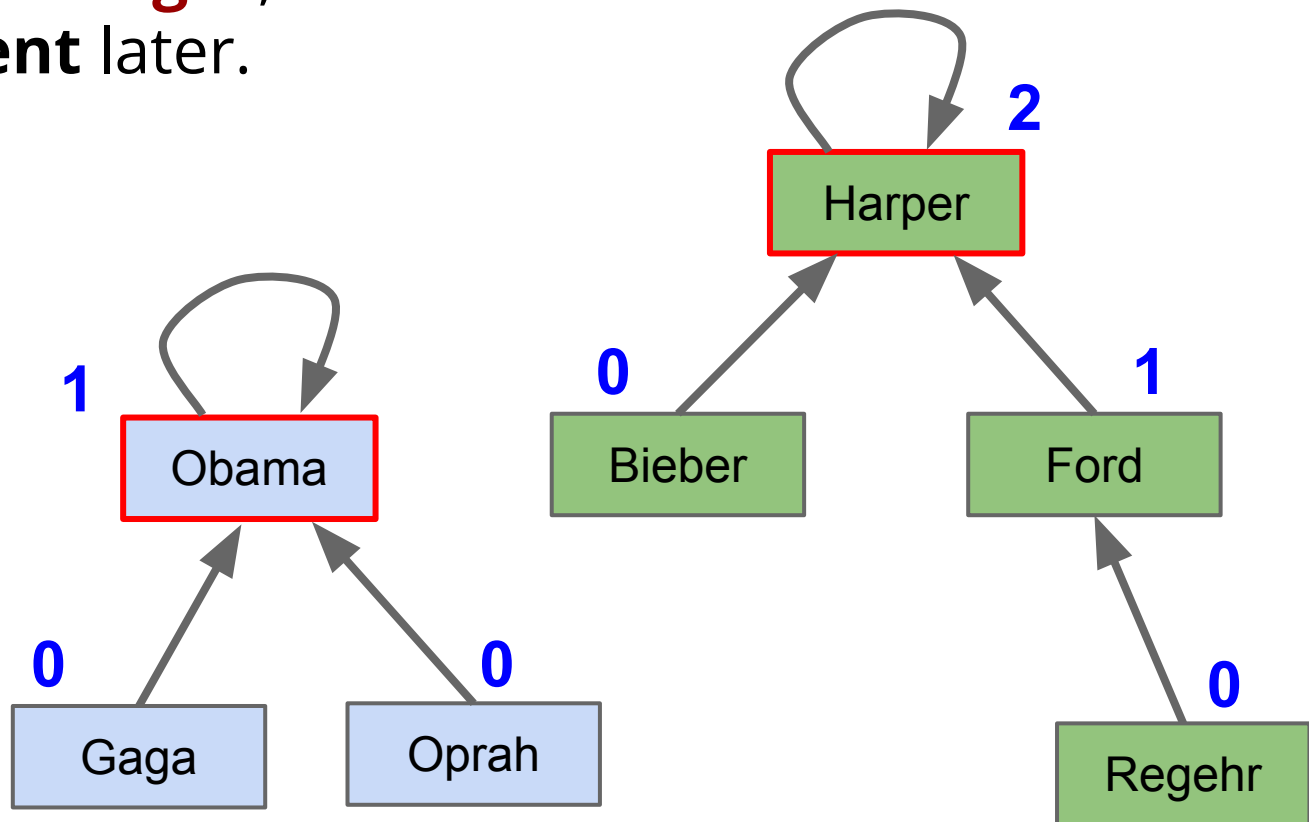
Intuition

- FindSet takes **$O(h)$** , so the **height** of tree matters
- To keep the unioned tree's height small, we should let the **taller** tree's root be the root of the unioned tree



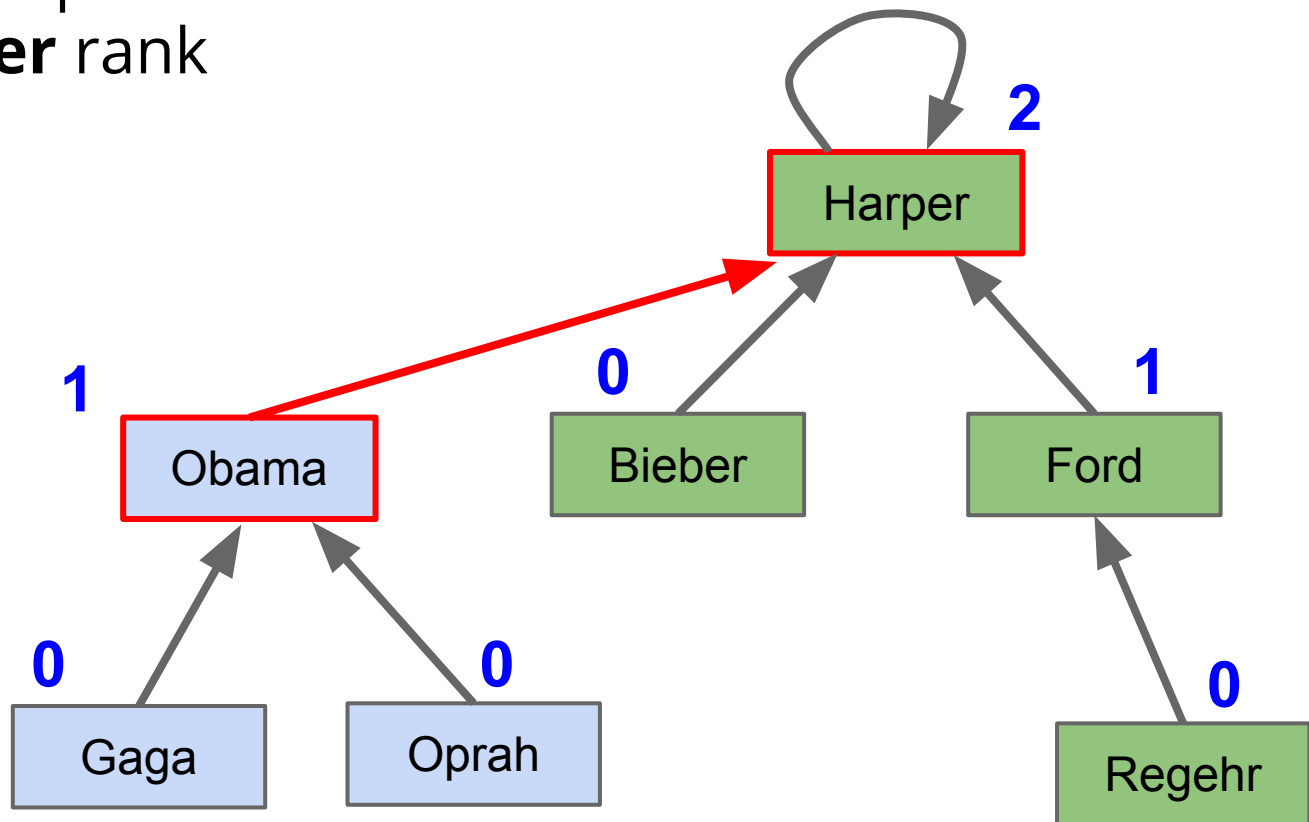
Each node keeps a **rank**

For now, a node's **rank** is the same as its **height**, but it will be **different** later.



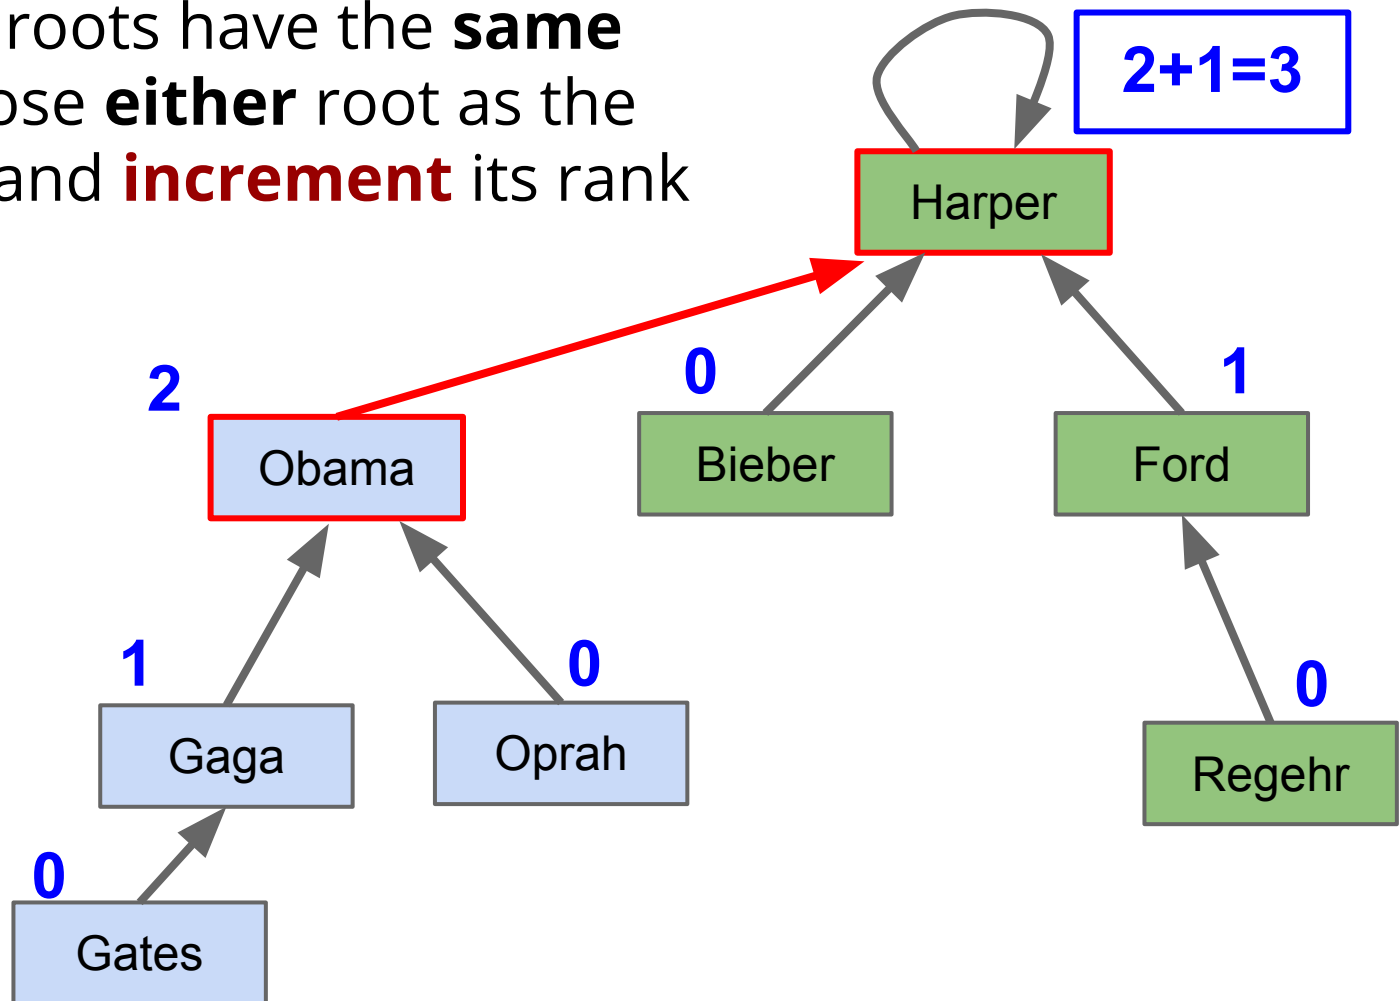
Each node keeps a **rank**

When **Union**, let the root with **lower** rank point to the root with **higher** rank



Each node keeps a **rank**

If the two roots have the **same** rank, choose **either** root as the new root and **increment** its rank



Benchmarking: runtime

It can be proven that, a tree of n nodes formed by **union-by-rank** has height at most $\lceil \log n \rceil$, which means **FindSet** takes $O(\log n)$

So for a sequence of $m/4$ MakeSets, $m/4 - 1$ Union, $m/2 + 1$ FindSet operations, the total cost is $O(m \log m)$

Rank of a tree with **n** nodes is at most **log n**,
i.e., **$r(n) \leq \log n$**

Proof:

Equivalently, prove **$n(r) \geq 2^r$**

Use **induction** on **r**

Base step: if $r = 0$ (single node), $n(0) = 1$, TRUE

Inductive step: assume $n(r) \geq 2^r$


→ a tree with root rank $r+1$ is a result of unioning two trees with root rank r , so

→ **$n(r+1) = n(r) + n(r) \geq 2 \times 2^r = 2^{(r+1)}$**

→ Done.

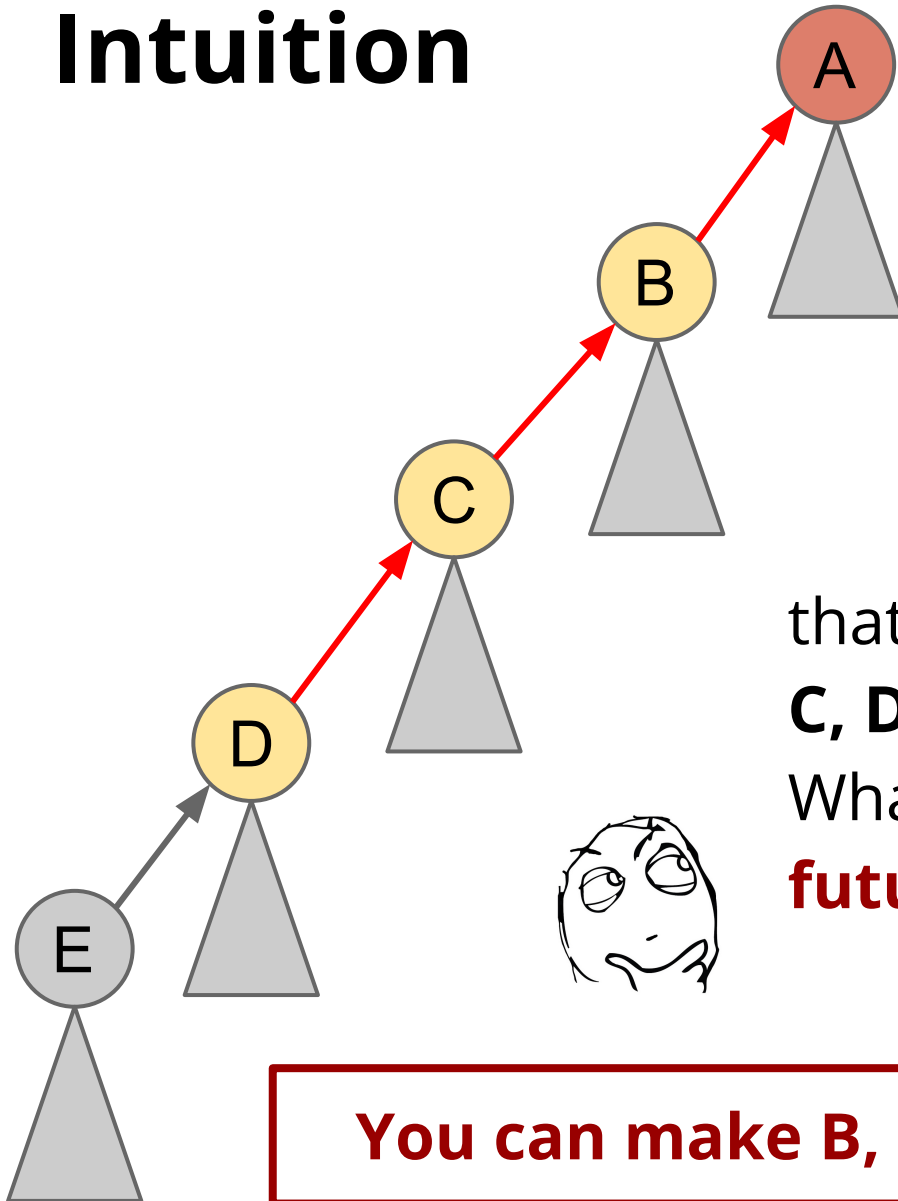
Trees with path compression

Intuition



Now I do a
FindSet(D)

Intuition




Now I do a
FindSet(D)

On the way of finding **A**, you visit **D**, **C**, **B** and **A**.


that is, now you have access to **B**, **C**, **D** and the root **A**.

What **nice** things can you do for **future FindSet** operations?



You can make B, C and D super close to A!

Make B, C and D
directly point to A



In other words, the path $D \rightarrow C \rightarrow B \rightarrow A$ is
“compressed”.

Extra cost to FindSet: at most **twice** the cost, so does not affect the order of complexity

Benchmark: runtime

Can be prove: for a sequence of operations with **n** MakeSet (so at most **n-1** Union), and **k** FindSet, the worst-case total cost of the sequence is in

$$\Theta \left(n + k \cdot \left(1 + \log_{2+\frac{k}{n}} n \right) \right)$$

So for a sequence of **m/4** MakeSets, **m/4 - 1** Union, **m/2 + 1** FindSet, the worst-case total cost is in **$\Theta(m \log m)$**

Ways of implementing Disjoint Sets

1. Circularly-linked lists $\Theta(m^2)$
2. Linked lists with extra pointer $\Theta(m^2)$
3. Linked lists with extra pointer
and with union-by-weight $\Theta(m \log m)$
4. Trees $\Theta(m^2)$
5. Trees with union-by-rank $\Theta(m \log m)$
6. Trees with path-compression $\Theta(m \log m)$

Benchmark:

Worst-case
total cost of a
sequence of m
operations
(MakeSet or
FindSet or Union)

Can we do better than $\Theta(m \log m)$?

U. B. R.

P. C.



Trees with union-by-rank and path compression

How to **combine** union-by-rank and path compression?

- **Path compression** happens in the **FindSet** operation
- **Union-by-rank** happens in the **Union** operation (outside **FindSet**)
- So they don't really interfere with each other, simply use them both!

Pseudocodes

Complete code using both union-by-rank and path compression

MakeSet(x):

$x.p \leftarrow x$

$x.rank \leftarrow 0$

FindSet(x):

if $x \neq x.p$: **# if not root**

$x.p \leftarrow \text{FindSet}(x.p)$

return $x.p$

Union(x, y):

$\text{Link}(\text{FindSet}(x), \backslash$
 $\text{FindSet}(y))$

Link(x, y):

if $x.rank > y.rank$:

$y.p \leftarrow x$

else:

$x.p \leftarrow y$

if $x.rank = y.rank$:

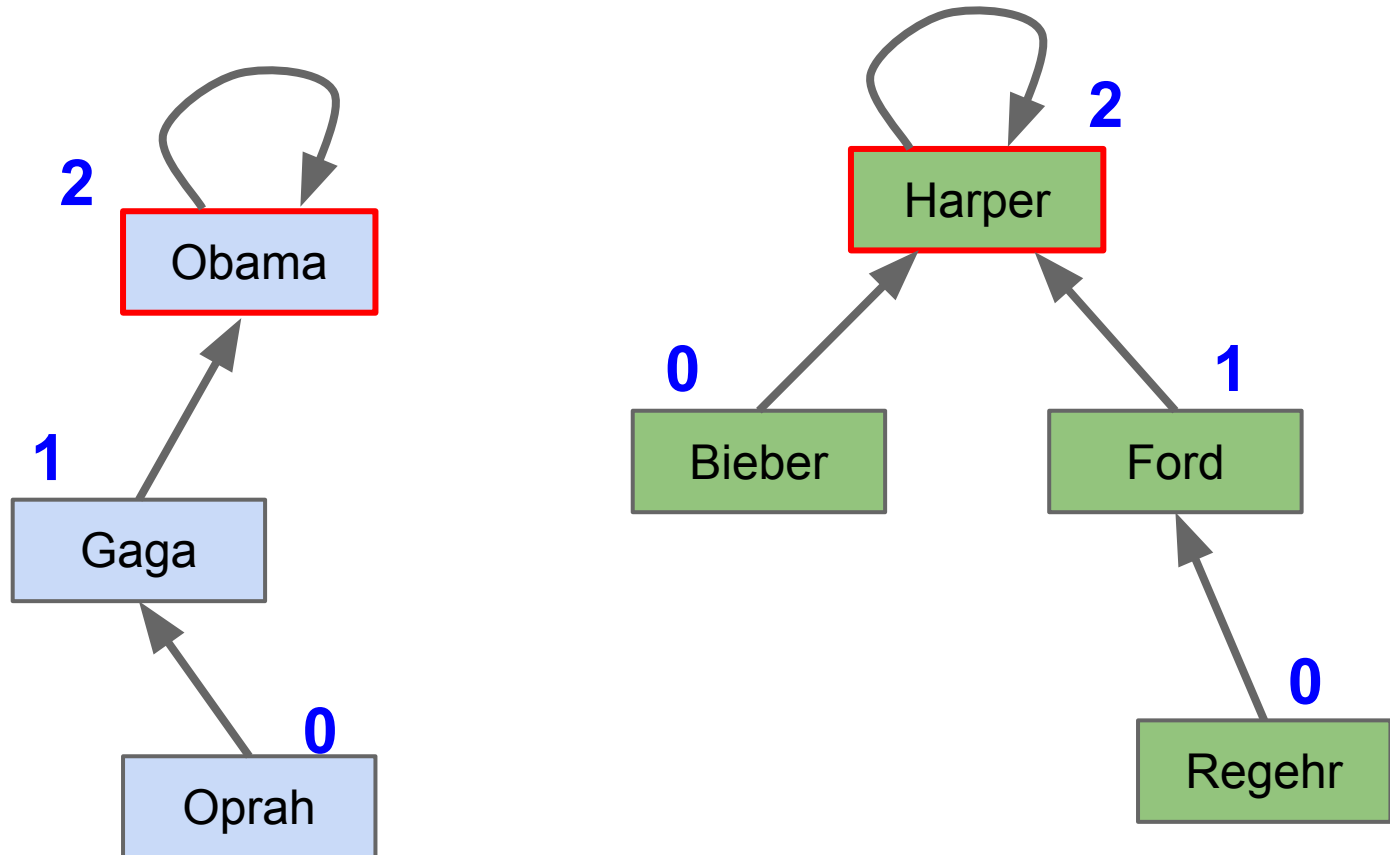
$y.rank += 1$



IT'S SO BEAUTIFUL
memegenerator

Exercise

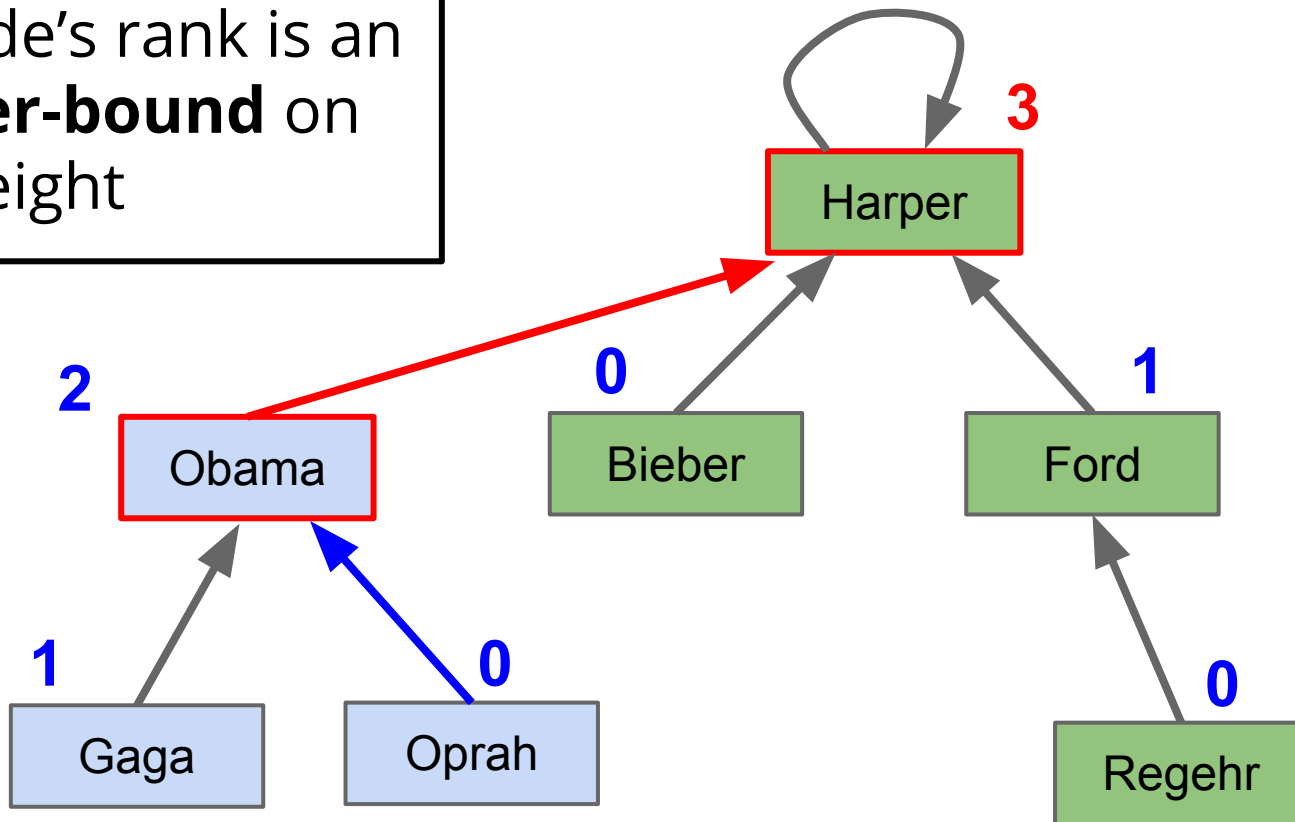
Draw the result after **Union(Oprah, Ford)**.
using both union-by-rank and path compression



Note: **rank** \neq **height**

because path compression does NOT maintain height info

a node's rank is an **upper-bound** on its height



Benchmark: runtime

Can be proven: for a sequence of **m** operations with **n** MakeSet (so at most **n-1** Union), worst-case total cost of the sequence is in

$$\mathcal{O}(m \cdot \alpha(n))$$

where **$\alpha(n)$** is the **inverse Ackerman function**, which grows really, really, really slowly.

In fact, **$\alpha(10^{80}) < 4$** , so we can basically treat it as **const**.

So the total cost of the sequence of **m operations is now improved to roughly **$\mathcal{O}(m)$****

Summary

1. Circularly-linked lists $\Theta(m^2)$
2. Linked lists with extra pointer $\Theta(m^2)$
3. Linked lists with extra pointer
and with union-by-weight $\Theta(m \log m)$
4. Trees $\Theta(m^2)$
5. Trees with union-by-rank $\Theta(m \log m)$
6. Trees with path compression $\Theta(m \log m)$
7. Trees with union-by-rank and
path compression $\approx O(m)$

Next week

- Lower bounds
- Review for final exam

<http://uoft.me/course-evals>