

Version Control

Let there be collaboration

The Challenge

Modern systems are too large for one person to build, but collaboration is difficult!

- Coordinating tasks
- Communicating quickly and efficiently
- Keeping everybody up to date, all the time
- As the team grows, so does the organizational overhead

Collaboration in Software

Collaboration is a general term, and can mean many different things:

- Sharing codebase
- Managing tasks & schedules
 - Keeping track of what needs to get done, by who and when
- Coordinating with "outsiders"
 - Other departments, such as marketing, design or business
 - Business partners' software team
- Spreading knowledge
- And more ...

Today's Focus

Sharing a codebase

- How does a team of software professionals work efficiently on a **single codebase**?
- What problems do they face?
- Which **tools and processes** do they use?
- How do these problems and tools change over time?

What do we mean by “sharing a codebase”?

Let's develop our requirements gradually ...

Fundamental requirement

Multiple developers should be able to edit a shared codebase

Very Simple Solution

- Shared folder + some conventions:
 - Read/write [source files](#) directly
 - Communicate via [email](#) (or [instant messaging](#)) to avoid stepping on each other's toes

Limitations

- No **backtracking**
 - What if someone accidentally breaks the code?
- No **traceability**
 - Changes since you last looked at the code?
- No **reliability**
 - Miss an email and you might accidentally overwrite your colleague's work

New Requirements

1. Shared codebase
2. History log, with the option to rollback to previous versions.

Simple Solution

- Use [Google Docs](#) (or any of its equivalents)
- Collaboration is built-in

Question: Where does this simple solution fail for software engineers?

Limitations

- Revision history at the level of individual files
 - Cannot (easily) rollback the whole codebase to a point in time
 - E.g.: `foo.h` and `foo.c` often change together
- No control over revision granularity
 - Auto-save is usually not ideal for coding
- No commit messages to go along with revisions
- Editor not suitable for coding

New Requirements

1. Shared codebase
2. History
3. Coding-specific requirements
 - Rollback the full codebase (not just individual files)
 - Commit when ready
 - Edit locally (using an IDE and/or text editor)
 - Define a snapshot of repo, namely a way of remembering version of every file at a point in time

You get the point ... Requirements evolve gradually over time

History Lesson

- Source Code Control System (**SCCS**) built in 1972
- Focus on **revision control for individual files**
 - Main revision history for each file
 - Ability to restore file(s) to a given revision
 - Ability to define snapshots using labels
 - Coordinate via lock → check-in → unlock cycles

History Lesson Continues

- `sccs` has some serious limitations
 - Runs on a single machine
 - As teams get larger, blocking others is impractical
- Version Control Systems have evolved, in order to solve these problems...

Client-Server VCS

- VCS software is installed on a **server**
- Code in an "**official repository**" on the server.
- Common workflow:
 - Checkout a copy of the code to your machine
 - Make changes **locally**
 - When the changes are ready (and tested), **commit** (i.e., save) changes back to the server
- VCS requires **conflicts to be resolved**, **before** they can be committed to the repository

Client-Server VCS

- Use **branches** to keep things manageable
 - Branch out to create **multiple versions of the codebase**
 - Compare differences between different branches
 - **Merge** changes from one branch to another
- Usage Examples:
 - Different branches for **different versions** of the product
 - **Feature** branches
 - **Experimental** branches

History Lesson

The Linux project quickly evolved into a very demanding project, in terms of version control.

- Many remote teams working on many different versions of the same codebase
- Different independent teams working (competing, even) to build same component. Conflict galore!
- Different time zones, languages and cultures
- All levels of competence and trustworthiness

Linux Project (before 2002)

- Different teams work independently on different versions, all based on some official repository
 - Periodically, “catch up” with the official repo
- When ready to contribute the code:
 - Send a patch file to Linus Torvalds for review
 - Linus (and other maintainers) fix problems and merge the patch at their discretion

What Do Patches Look Like?

```
diff -Naur base/foo.c changed/foo.c
--- base/foo.c 2015-09-13 10:51:35.000000000 -0400
+++ changed/foo.c 2015-09-13 10:51:17.000000000 -0400
@@ -1 +1,2 @@
 void f(){ /*implement f*/}
+void f2(){ /*implement new function*/ }
diff -Naur base/foo.h changed/foo.h
--- base/foo.h 2015-09-13 10:51:54.000000000 -0400
+++ changed/foo.h 2015-09-13 10:51:48.000000000 -0400
@@ -1,2 +1,3 @@
 //interface file
 void f();
+void f2(); //add this definition
```

Do you think staring at diff output is a lot of fun?

Applying patches

- Nasty work, but someone has to do it
 - Applying patches is tedious
- Nevertheless, the Linux team was onto something important
 - Workflow for many autonomous teams
 - Contributors package their changes, and send them as a single unit to the project maintainer(s)
 - **Goal:** Distribute responsibility/work among contributors, and prevent maintainers from becoming a bottleneck
- This workflow evolved into what we know as *pull request*

End Of History Lesson

- In 2002, the Linux project started using BitKeeper
 - Distributed Version Control System - A new type of VCS that tackled the collaborations issues that the Linux team (among others) was experiencing.
 - Proprietary (and *not* open-source) technology
- In 2005, the Linux project switched to Git
 - Git was built by Linus Torvalds as a replacement for BitKeeper
 - The reason for the switch ... licensing disagreement

New Requirements

1. Shared codebase
2. History
3. Coders' requirements
4. Better support for Linux style collaboration:
 - Independent versions (i.e., different repos), based on an official repo
 - Easily catch up with the official repo
 - Pull request

Solution - Distributed VCS

- No single “official” repository
- Repositories (with their complete history) can be **cloned** at any point in time
- Changes (aka commits) can be **pushed/pulled between repositories**
- Every commit is a **snapshot of the filesystem**
- A repo is **graph of commits**

Advantages Of Distributed VCS (DVCS)

- Allow for ***Pull Requests***

- A better way to request someone to pull a given set of changes into their codebase
- Easier for project maintainers to **distribute duties** (e.g., resolving conflicts) to team members
- Essentially, **collaboration** is built into the version control system

Other Advantages Of DVCS

- Better *branching* and *merging*
 - Allow you to **maintain multiple versions of the same codebase**
 - Much easier when you think of **snapshots**, instead of changes to individual files.
- Work **offline** and push changes later
 - E.g.: Working on an airplane
- Control over commit history, which documents the evolution of your system

Open Source Software

- DVCS has had a significant positive impact on Open-Source Software (OSS)
- In many OSS projects:
 - Team members are [remote](#)
 - Different time-zones and languages
 - May be complete [strangers](#) ([cannot assume trust](#))
- DVCS allowed us to build better tools for open-source software development

History Lessons: Summary

- The landscape is constantly changing
 - More demanding requirements lead to more powerful tools, which lead to even more demanding requirements ...
 - E.g.: SCCS → CVS → SVN → Git → ?
- But one concept always stays the same - Professionals use tools to become more productive

Version Control in CSC301

- Git
 - Industry standard DVCS
 - Open-source
- GitHub
 - Hosting service for Git repositories
 - Web-based toolset for code/project management.
 - Free for public projects
 - Industry standard for OSS development
 - Atlassian Bitbucket, is a similar tool

Enough motivation, let's see Git ...