

CSC301

Asynchronous patterns & Object pools

Announcements

- A4
 - Due **March 6, 23:59**
 - Marks out later this week (tentatively)
 - In the `auto-marker` branch of your A4 repo (the original one, not your fork)
 - Please let us know if there is a mistake
- Team deliverable 2 is due **March 9, 23:59**

Multithreading

- Multiple threads of execution running “simultaneously”
 - Well ... not exactly simultaneously, but sharing the CPU makes it seem like it
 - Java also supports creating new processes, but threads are much more common
- Multithreading is essential to most modern applications
 - Servers handling multiple simultaneous requests
 - GUI applications (GUI updates and background computations happen on separate threads)
 - Applications using 3rd party, remote API's
(applications cannot hang while waiting for network response)

Some Benefits of Multithreading

- Improved **resource utilization**
- Increased **responsiveness**
 - Chunks of work can be handed over to worker threads - e.g., servers handling requests, GUI
- Efficient use of **multiprocessors**
- Improving **program structure**
 - Simplified code, less need to coordinate concurrent tasks

Multithreading

- Multi-threaded programming is traditionally considered difficult
 - Thread interaction can be difficult
 - Easy to get things wrong
 - For example: [deadlock](#), [race condition](#), [starvation](#), etc.
 - Harder to reason about and debug
 - (mainly because the order of execution is [nondeterministic](#))
 - Errors difficult to detect, reproduce, and fix

Java & Multithreading

- The API's related to threads (and concurrency) evolved over the years
- Initially, there was just Thread
 - Either extend it (i.e., create a subclass) or construct it with a Runnable
 - Once you create a Thread instance, you can start it
 - Fairly low-level methods allow synchronization and inspection of threads

Asynchronous Programming

- Say we have a (fairly long-running) task that returns a result
 - E.g.: Authenticate using a remote service (e.g., Google, Facebook, GitHub, etc.)
- As a programmer, I want to:
 - Run the task in a background thread, while doing other things in the main thread
 - Do something with the result of the task, whenever it's ready
 - Handle any error that the task might throw
 - Specify a timeout, so that my program is not stuck waiting for the task, in case it runs indefinitely

Asynchronous Callbacks

- Instead of waiting for a function to return a result, **tell it in advance what you want to do with its result**
- That is:
 - When calling the function, pass a **callback** (the callback is just another function)
 - The function runs in a **separate thread**
 - Once the function is done, the **callback is called** (with function's result as its argument(s))

Synchronous vs. Asynchronous

Let's see the difference between synchronous and asynchronous code.
(*Note:* This example uses JavaScript)

```
try {  
  
    data = syncFunc(someArg);  
    // Do something with the data ...  
  
} catch (error) {  
    console.log(error);  
}
```

```
asyncFunc(someArg,  
  
    function(data) {  
        // Do something with the data ...  
    },  
  
    function(error) {  
        console.log(error);  
    }  
);
```

Callbacks in JavaScript

- Functions might take a long time to complete
 - `var photo = downloadPhoto('http://coolcats.com/cat.gif')`
`// photo is 'undefined'!`
- Don't want to block while the photo is downloading
 - Store the code that should run after the download is complete in a function (a [callback](http://callbackhell.com/))

```
downloadPhoto('http://coolcats.com/cat.gif', handlePhoto)

function handlePhoto (error, photo) {
  if (error) console.error('Download error!', error)
  else console.log('Download finished', photo)
}

console.log('Download started')
```

Asynchronous Programming

More features/requirements:

- Chain tasks together
 - When a task is done, pass its result to the next task (and run the next task)
 - More complex “plumbing”, for example:
 - When a task is done, start three different tasks (in parallel)
 - Wait for multiple tasks to complete (all of them), and only then run the next task
 - Wait for the first of multiple tasks to complete, and pass its result to the next task
 - Pass exception down the chain, handle it and (possibly) continue to chain additional tasks
- Specify which threads to use for running background tasks
 - More control

Chaining Asynchronous Calls

```
asyncFunc1(someArg,  
  function(result1){  
    var x = doSomethingWith(result1);  
    asyncFunc2(x,  
      function(result2){  
        doSomethingWith(result2);  
      },  
      function(error){  
        console.log("f2 error: " + error);  
      }  
    );  
  },  
  function(error){  
    console.log("f1 error: " + error);  
  }  
);
```

In the example on the left, we chain two asynchronous functions.

Essentially, we are programming the following logic (plus error handling):

Call `asyncFunc1` and, when its result is ready, call `asyncFunc2`.

Chaining Asynchronous Calls

As our chains get longer (and, possibly, more complex), the code becomes very hard to manage.

For example:

```
getData(function(a){
  getData(a, function(b){
    getData(b, function(c){
      getData(c, function(d){
        getData(d, function(e){
          ...
        });
      });
    });
  });
});
```

Note: This pattern is known as [Callback Hell](#)

Drawbacks of multithreading

- Harder to trace and debug
- Exponential growth of complexity when multiple threads are running
- Higher probability of unexpected results (i.e., nondeterminism)
- Unintended dependencies (e.g., exception handling)

Promises (Futures)

- Asynchronous programming fits well with many of the tasks modern programmers deal with (because of the nature of the Internet)
- That being said, it doesn't always fit well with the syntax that modern programmers are used to
 - Chaining multiple methods can result in code that is very hard to read
 - Exceptions need to be handled more carefully
- The ***promise*** design pattern to the rescue...

Without promises

```
asyncFunc1(someArg,  
  function(result1){  
    var x = doSomethingWith(result1);  
    asyncFunc2(x,  
      function(result2){  
        doSomethingWith(result2);  
      },  
      function(error){  
        console.log("f2 error: " + error);  
      })  
    ),  
  function(error){  
    console.log("f1 error: " + error);  
  })  
);
```

vs.

With promises

```
asyncFunc1(someArg)  
  .then(function(result1){  
    return doSomethingWith(result1);  
  })  
  .then(function(result2){  
    doSomethingWith(result2);  
  })  
  .catch(function(error){  
    console.log("Caught " + error)  
  });
```


Promise

- **Promise** (aka *future*, *delay* or *deferred*) is an object that represents a **result that is initially unknown**
- Different languages use different constructs and terminology (e.g., Java uses CompletableFuture), but the basic idea is the same:
 - **Decouple a computation from the value it produces**
 - Compose computations (i.e., functions) together (before you run/execute them)
 - A nice tutorial on using CompletableFuture (and another one)
 - Can promises fail to complete successfully? How do we handle that?
- Where can we leverage the asynchronous nature of promises/futures?
 - Computationally intensive processes (e.g., scientific calculations)
 - Manipulating large data structures
 - Remote method calls (downloading files, web services, etc.)

Java Threads - More Control

- Threads are resources (just like memory or CPU time)
- Sometimes we want more control over how resources are managed

For example:

- Parallelism
- Reusing threads (instead of creating new ones)
- Managing cache (to optimize context switching)
- etc.

Resource Pool

- Object Pool is a common design pattern (aka Resource Pool)
 - As the name suggests, we create an object that is responsible for managing a pool of resources
 - The resource pool manages the **lifecycle** of resources (e.g., construction, destruction, etc.)
 - Two of the most common examples for resources that can benefit from pooling:
 - Threads
 - Database connections
 - ScheduledThreadPoolExecutor is an example of a thread pool that comes built-in with Java
 - Here is really nice tutorial you might find useful

Code Example

- This commit history of this short [code example](#) shows a few of the concepts discussed in this lecture.

Example - Context

- An application needs to process hundreds of thousands of records and update them
- Process is a multi-step transaction and the order of execution is important
- Records are grouped by “Accounts” that are independent of each other
- Records may need other resources that are shared among them
- Speed is not important, but success and correctness are critical

Example - Initial Solution

- Creating a new thread for every account (i.e., pool)
- Chaining multiple steps using promises (futures)
- Shared resources are locked when they are needed

Example - Problems

- Unpredictable results
 - Dependent on the order of execution of threads for each account
 - Deadlocks
- Complicated process to handle unexpected results
- High resource consumption

Example - Final Solution

- Simplified multiple steps into two
 - Minimized the number of chains
- Made every process completely independent
 - Architected around shared resources
- Grouped all DML (Data Manipulation Language) operations into one final “clean-up” process
 - Created one shared final cleanup (similar to a shutdown hook)

Summary

- We use multithreading to perform multiple tasks simultaneously
- We can use *thread-pools* to manage our threads.

Common example:

- The UI thread (a pool with a single thread)
- Worker threads (use for long-running background tasks)
- When working with asynchronous code:
 - Use ***promises*** to chain functions together (i.e., create pipelines) and keep the code easily maintainable
 - When executing an asynchronous pipeline, use a thread-pool for long-running, background tasks (e.g., network request, or reading from disk)