

CSC301

Distributed Applications

Midterm Information

- Next week - **March 19 @ 18:10 OR March 20 @ 12:10**
 - Duration: 90 minutes
 - **Attend the section where you are officially registered!**
- Questions
 - Git/GitHub
 - Given a scenario, write a series of commands to accomplish a certain task
 - Java Programming
 - Single domain
 - Given: interfaces, helper functions, stubs, test cases
 - To-Do: complete implementations of classes/methods using design patterns studied in the course
 - **No Aids**

Guest Lectures

- Planned for **Monday/Tuesday March 26/27**
- Speakers from the industry (aiming for 2)
- Open discussion format
 - Their experience working in development and managing development projects
 - Tools, processes, techniques they use
 - How development is happening in their organizations - e.g., a feature journey
 - Hiring practices
- Other topics?

Today's goal

The goal of this lecture is to connect (some of) the dots between

- What you see in your programming courses, and
- The technology you use in your daily life.

Working With Multiple Processes

- Up until now we have been discussing objects that are in the same Java Virtual Machine (same process, same machine).
- However, almost any significant modern application spans several processes and/or machines.
- How do we communicate between processes?

Network communication

Remember CSC209? How did you communicate between processes?

- **Pipes & local files** - between processes on the same machine.
- **Sockets** - Allow you to communicate between machines.

Network communication

- Sockets

- Essentially a “pipe between machines”
- Two processes open a socket between them, and read/write data (aka send/receive messages).

- Protocols

- Sometimes (usually) the data needs context.
- The format of the data on the wire.

What does `foo.f()` mean?

- When we see
 - `foo.f()`;
- We read
 - “*calling*” a function `f` on object `foo`
- In some OO languages, we say:
 - “*Sending the message*” `f` to object `foo`

Sometimes, it’s easier to imagine “sending a message” between machines than “calling a function” across machines.

Distributed Objects

- How should we think about an app that sends data between machines?
 - How about a network of “distributed” objects that “send messages” to each other?
- How should we think of the messages?
 - The receiver (`self` or `this`, depending on language)
 - The name of the method
 - Arguments to the method

Remote Procedure Call (RPC)

- A **remote procedure call** is a method call where the callee is on one machine and the caller is on another.
 - Developers don't explicitly code the details of these interactions - the code is generally very similar to a local procedure call.
- How do the arguments travel from callee to caller?

Remote Procedure Call (RPC)



- A **remote procedure call** is a method call where the callee is on one machine and the caller is on another.
- How do the arguments travel from callee to caller?
 - **Serialization/deserialization** to convert arguments (objects) to/from data (bytes)
 - Send/receive the resulting data through a **socket**!

Examples

- Sending bytes over a socket
- Sending a Java-serialized object over the socket
- Java RMI - Calling a method on a remote object, as if it was local.
 - See this [hello world example](#), where the RMI library
 - Allows you to bind/lookup objects by name
 - Provides you with proxy objects that abstract away all communication details
 - RMI depends on agreement/convention between the participants (e.g., client and server)
 - *Note:* RMI is a bit old and is not so popular these days.
 - But the same concepts/challenges apply to current libraries/technologies as well.

Java RMI

- Remote Method Invocation - OOP version of RPC
 - Simplifies interaction between remote Java Applications
 - Caller delegates request to its stub (proxy object), a local object representing a remote object
 - **Stub** passes caller arguments over the network to the server skeleton
 - **Skeleton** passes data to the called object, waits for response and returns the results
 - Stub/skeleton responsible for **marshalling/unmarshalling messages**



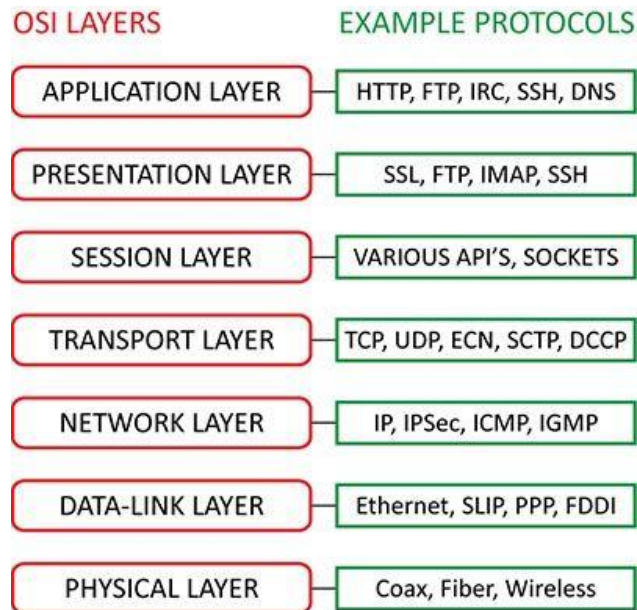
[Image credit: Wikipedia]

Distributed Applications

- Processes on different machines need to communicate
- At the lowest level, use **sockets**
 - Send/receive bytes
 - Control every bit that goes on the wire
- Q: How to get from pushing bytes through a socket to the Internet as we know it?

Protocols

- *Hierarchy* of protocols on top of sockets
 - Agreed-upon standards/conventions
 - Enable communication
 - Each protocol *abstracts a level of details*
- Participants (i.e., processes) know
 - When they should send data
 - What data to send
 - When they should wait for incoming data
 - What incoming data is expected to look like



Basic Communication Protocol

- IP - Network Layer
 - The principal communication protocol of the Internet
 - All about routing - “Getting data from A to B”
- TCP - Transport Layer
 - Transport-layer protocol that complements IP
 - **Reliable** data transfer using error detection
 - Alternative - UDP (User Datagram Protocol), unreliable replacement for TCP

Basic Communication Protocol

- TCP/IP is a *transport-layer* protocol
 - Low-level protocol(s), on top of sockets
 - Not much semantics, all about transferring data reliably across the wire.
- If we want to build the Internet, we need ...
 - To think a bit more high-level
 - To abstract away details of moving bytes on a wire
- We need an *application-layer* protocol

Modern Web Applications

- Java-only protocols are no longer an option
- For example, think of Facebook's system...
 - iOS app (objective C/Swift)
 - Android App (Java)
 - Web client (JavaScript)
 - Web server (PHP)
 - And many (many!) more internal pieces that make up their extremely complex and large system.

Modern Web Applications

- Need a communication protocol that is
 - Application-layer
 - Allow us to focus on building applications, not arranging bits on the wire.
 - Language-agnostic
 - Exchange data between applications written in different programming languages
 - Adopted by the industry as a standard!

HTTP

- Application layer protocol
 - On top of reliable transport layer (e.g., TCP/IP).
- Client-server protocol
 - Client sends a request
 - Server sends back a response
- One of the most commonly-used protocols on the web

HTTP Request

- Request line
 - Indicates which resource the client is requesting
 - E.g.: `GET /index.html HTTP/1.1`
- Request headers
 - Contain metadata
 - E.g.: `Accept: application/json`
- An empty line
- Optional message body (i.e., data)

HTTP Response

- Response line
 - Includes a status code
 - E.g.: `HTTP/1.1 200 OK`
- Response headers
 - Contain metadata
 - E.g.: `Content-Length: 348`
- An empty line
- Optional message body (i.e., data)

Live Demo

1. Open your browser at <http://example.com>
2. `telnet www.example.com 80`, and type

```
GET /index.html HTTP/1.1  
Host: www.example.com
```

(With an empty line at the end)

Example Response

Request

```
telnet www.example.com 80
Trying 93.184.216.34...
Connected to www.example.com.
Escape character is '^]'.
GET /index.html HTTP/1.1
Host: www.example.com
```

Metadata

```
HTTP/1.1 200 OK
Cache-Control: max-age=604800
Content-Type: text/html
Date: Mon, 12 Mar 2018 14:57:41 GMT
Etag: "1541025663+ident"
Expires: Mon, 19 Mar 2018 14:57:41 GMT
Last-Modified: Fri, 09 Aug 2013 23:54:35 GMT
Server: ECS (lga/1383)
Vary: Accept-Encoding
X-Cache: HIT
Content-Length: 1270
```

Response Body

```
<!doctype html>
<html>
<head>
  <title>Example Domain</title>

  <meta charset="utf-8" />
  <meta http-equiv="Content-type" content="text/html; charset=utf-8" />
  <meta name="viewport" content="width=device-width, initial-scale=1" />
  <style type="text/css">
    body {
      background-color: #f0f0f2;
      margin: 0;
      padding: 0;
      font-family: "Open Sans", "Helvetica Neue", Helvetica, Arial, sans-serif;
    }
    div {
```

• • • •

Live Demo - Cont'd

- In both cases, you sent HTTP request to a server, and got a response.
- Let's look at the telnet window
 - Save the response body (without the headers) to a local file, test.html.
 - And open the file in your web browser.

Another Demo

- Every modern browser has *developer tools*
 - Allow you to inspect what the browser is doing under the hood
 - Different browsers = different menu/shortcut
- Let's try something ...
 - Open the developer tools in your browser
 - Go to <http://google.com>

Another Demo - Cont'd

- We can see all HTTP requests:



- We only asked for one website, how come there are 25 requests?

Another Demo - Cont'd

- In your browser, keep the *developer tools* open, and start typing a search query in Google.
 - The browser sends HTTP request(s),
 - Gets a response containing JSON data,
 - And updates the screen (i.e., display auto-complete options)
- Notice ... All network communication happened in the background

Sync vs. Async

- Synchronous flow
 - Call a function, wait for the result
- Asynchronous flow
 - Call a function
 - Go do something else (without waiting for the result)
 - When the function returns (at some point), do something useful with its result.

Very Quick “History Lesson”

- Mid-90's
 - Browsers were working synchronously
 - User-experience was inferior (in today's standards)
- Late-90's
 - Browser vendors gradually added async support
- Early 2000's
 - Developers realized the potential
 - Emerging standards, conventions, tools and hypes.

AJAX

- Asynchronous JavaScript And XML
 - Decouples data interchange and presentation layers, allowing for dynamic content change (no need to reload the entire page)
- Huge hype around 2004
 - GMail and Yahoo Mail were some of the early notable web apps to offer desktop-like UX
- In time, JSON replaced XML as the standard serialization format
 - Less data to transfer
 - Easier to parse in web browsers (that have an extremely optimized JavaScript engine)

Async Function Call - Review

- When calling f , provide a callback
 - Like Lambda expressions or function pointers
 - Possibly multiple callbacks (e.g. success & error)
- Run f in a background thread
 - Other threads don't have to wait for f
 - Extremely useful if f makes a network request
- When f returns a result, call the callback function (with the result as an argument)

Software As A Service

- Instead of providing users with software that runs on their machine, provide them with a service that runs on yours.
 - Easier to maintain
 - For many businesses, solves piracy issues
- Client talks to your software over the web
- More standards and conventions \Rightarrow More powerful tools & libraries

SaaS

- This is how modern web/mobile apps work
- Use remote service(s) to
 - Show weather info
 - Manage comments on articles
 - Display ads
 - etc.

SaaS

- Each service has its own API
 - Similar to interfaces in Java
- The API specifies
 - Format
E.g.: JSON over HTTP
 - Available resources
E.g.: GET /weather?city=Toronto,CA HTTP/1.1
 - Expected request/response data
E.g.: {"text" : "hot", "degrees" : "33c"}

Summary

- Sockets
 - Let us transfer bytes *reliably* between machines
- Communication protocols
 - Abstract low-level communication details, and used for communicating between applications
- Serialization
 - Convert in-memory objects to/from data that can be sent to/from different machines
- Threads
 - Allow us to build asynchronous flow
 - Async flow enables us to use rich web-services while providing quality UX
 - **Lambdas** (i.e., functions) and **promises** make async code more elegant and maintainable
- Conventions & Standards
 - E.g.: By following **REST**, one can **auto-generate server and client code from documentation**.