

More on Testing and Large Scale Web Apps

Testing

- Functionality Tests
 - Unit tests: E.g. Mocha
 - Integration tests
 - End-to-end - E.g. Selenium
 - HTML CSS validation
 - forms and form validation
 - cookies - test for correct operation and deletion

Testing

- Usability testing
 - need real people to test your application
 - navigation
 - consistency
 - content

Testing

- Interface testing
 - application - requests sent to server and results displayed
 - web server - handles all requests correctly
 - database server - queried give expected results
 - when connection between layers fails, appropriate error messages
- Database testing
 - integrity
 - response time
 - data retrieved is shown accurately

Testing

- Compatibility testing
 - works on all major browsers, including older versions
- Performance testing
 - load test - normal and peak
 - stress test - push beyond peak
 - crash recovery
- Security testing
 - regularly audit!

Data-driven decision making

(Kate Hudson - Mozilla)

- In the world of continuous deployment, how and when are decisions made?

Healthy software?

- Doesn't crash (very often)?
- Trustable data
- Preserves data
- Easy to use
- Performs well (?)
- Handles errors gracefully

Performs well?

Goal: Minimize user-perceived delay

- Reduce time to fetch a resource
 - server processing time
 - browser caching
 - HTTP pipelining
 - minimize JS resources and use shared libraries
- Optimize rendering speed
 - layout
 - CSS
 - Javascript optimizations
- Make loading time appear shorter
 - loading indicators
 - deferred loading

Monitoring

- Detecting when software is not healthy
- Monitoring software: CPU, memory, #crashes
- Look at HTTP response codes: 500, 400

Roll-out

- Server-side: Canary release pattern
 - update one server, check monitoring
- Client-side:
 - staged roll-out - to a small group of users
 - must be able to split up users
 - some percentage get new code
 - by single country or language (not en-us)
 - opt-in release channels (beta channel)
 - can collect more data, bug reports, feedback

Firefox

- Different versions depending on risk
 - beta, nightly, aurora
 - users expect some instability
- Prefs - [about:config](#) in Firefox to turn on/off experimental features
- Test Pilot - add on where users can try out new features

Measure User Behaviour

- User engagement:
 - track clicks
 - retention

AB testing

- Control group and experimental cohort - users unaware of participation
- Change some aspect of the application:
 - order of operations
 - amount of information shown at one time
 - placement of buttons
 - text of buttons
 - choice of images

AB testing


- Important to run one experiment at a time
- Run long enough so that statistical significance
 - aim for 90-95% certainty
 - need a lot of users to do this
- Results must be correlated with right metrics
 - sometimes conclude that neither is more effective
 - sometimes see unexpected effects - e.g. a huge drop in engagement
 - novelty effect

Ethical considerations


- What user data are you collecting?
- Impact on user performance?
 - E.g. Research question: Would a degradation in performance lead to lower user retention?
 - Artificially reduce introduce latency for some users.
 - Is this ethical?

Large Scale

Standard Web App



Large Scale Web App



Scale-out

- Expand capacity by adding more instances
- Pros:
 - can scale to fill needs by adding and removing instances
 - fault tolerance
- Cons:
 - manage multiple instances and distribute work

Scale Out: Which server to send to?

- Browsers want to speak HTTP to a web server
 - Use load balancing to distribute incoming HTTP requests across many front- end web servers
- HTTP redirection:
 - Front-end machine accepts initial connections
 - Redirects them among an array of back-end machines
- DNS (Domain Name System) load balancing:
 - Specify multiple targets for a given name
 - Handles geographically distributed system
 - DNS servers rotate among those targets
 - How to handle sessions?

Load-balancing switch ("Layer 4-7 Switch")

- Special load balancer network switch
 - Incoming packets pass through load balancer switch between Internet and web servers
 - Load balancer directs TCP connection request to one of the many web servers
 - Load balancer will send all packets for that connection to the same server.
- In some cases the switches are smart enough to inspect session cookies, so that the same session always goes to the same server.
- Stateless servers make load balancing easier (different requests from the same user can be handled by different servers).
- Can select web server based on random or on load estimates

nginx

- Web server designed for scalability
- Load balancer
 - can handle SSL processing
 - application health checks (server fails)
 - session persistence
 - limits to mitigate DOS
 - bandwidth limiting

Scale-out assumptions

- Any server will do
 - Different requests from the same user can be handled by different servers
 - Requires database be shared across servers
- What about session state?
 - should be fast because it is accessed on every request
- Web sockets?
 - cannot load balance each request

Scalability and Sessions

- Pass session state around in cookies if possible
- Cache session state in memory, but store in Database
- Sticky sessions
 - Load balancing - hash on client IP
 - With TLS/SSL offloading can use session cookie

Scale-out storage

- Data sharding - spread database across instances
 - each piece of the database is called a shard
 - Tolerate failures (and improve performance) by replication
 - Increases complexity - Applications must place data across multiple databases
- A Router decides where to send the request(s) to.
- Key questions:
 - how to organize data across shards
 - by key, by tag, geolocation?
 - replication

Memcache

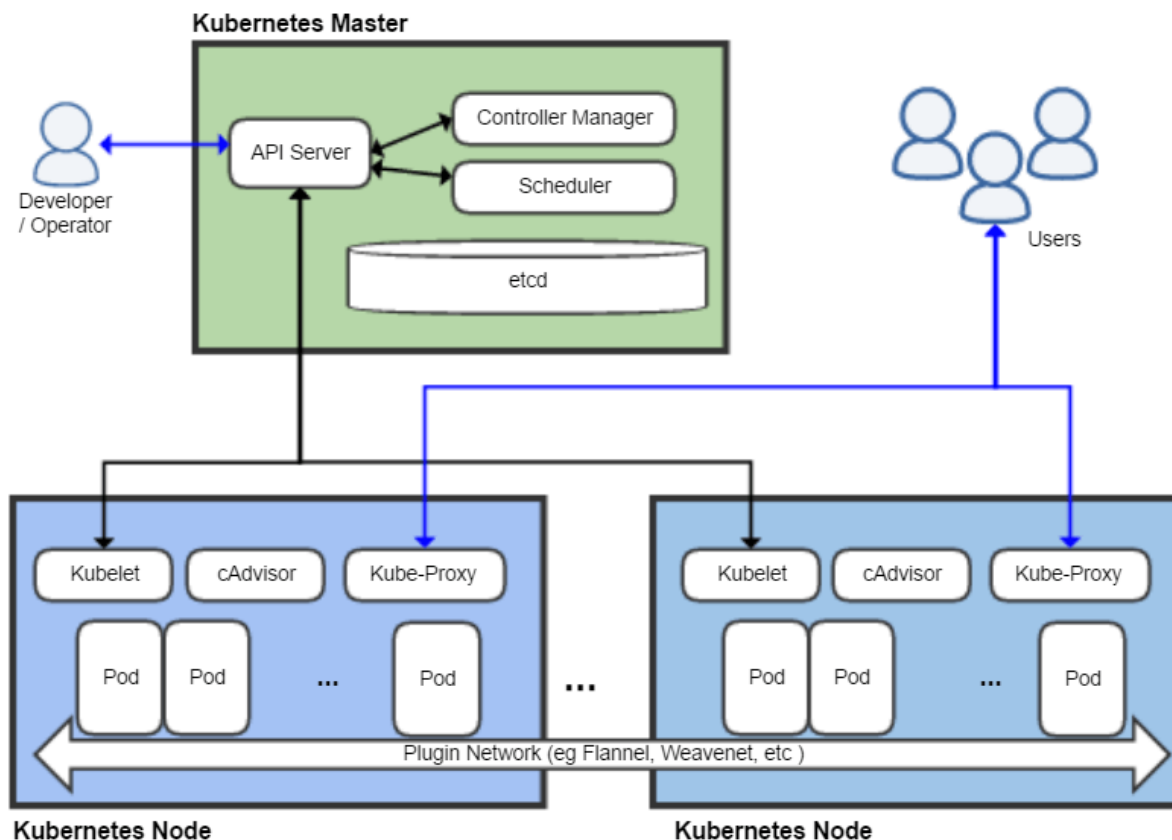
- Main memory caching system
- Key-value store (both keys and values are arbitrary blobs)
- Used to cache results of recent database queries
- Much faster than databases:
 - 500-microsecond access time, vs. 10's of milliseconds
 - Writes still go to database, so no performance improvement
 - Cache misses still hurt performance
 - Must manage consistency in software (flush memcache data when database is modified)

Scalability

- Building this architecture is hard!
- Need data centre expertise
- Figuring out the right number of components is hard
- Cloud computing:
 - allows for dynamic addition and removal of resources
 - outsources data centre management

Kubernetes

“Open source system for automating, deploying and managing containerized applications”



Deployment options

- ngrok
 - public URL for exposing local web server
 - relays traffic through ngrok process running on local machine
- heroku
 - public place to deploy web apps
 - supports wide range of technologies
- glitch.com
 - build web apps in the browser
- Firebase
 - back end as a service
 - real time database (web sockets to NoSQL DB)
- Others (?)